Abstract

The characterization of viscoelastic fluids presents significant challenges in both industrial processing and fundamental rheology. Accurately modeling non-Newtonian behavior requires non-trivial governing equations coupling momentum conservation with frame-indifferent constitutive laws. To investigate this rheological complexity, the four-roll mill provides a benchmark geometry: by modulating the relative rotation speeds of four cylinders (Ω and $-k\Omega$, with $k \in [-1, 1]$), the apparatus generates pure rotation, simple shear, pure planar elongation, and continuous mixed kinematic regimes in between, providing an ideal platform to probe viscoelastic response in these less ideal regimes.

Characterizing viscoelastic properties from experimental velocity data (such as Particle Image Velocimetry) constitutes an inverse problem that standard Computational Fluid Dynamics (CFD) is not well suited to solve. In fact, recovering unknown constitutive parameters via mesh-based solvers requires an expensive, iterative trial-and-error procedure involving repeated simulations.

This thesis addresses these challenges by developing a Physics-Informed Neural Network (PINN) framework specifically tailored for the four-roll mill—a benchmark geometry capable of generating diverse shear and extensional flow regimes.

By embedding the governing conservation laws and frame-indifferent constitutive equations directly into the neural network's loss objective, the proposed framework enables the simultaneous reconstruction of unmeasured latent fields, such as polymeric stress and hydrodynamic pressure, and the identification of unknown rheological parameters. The framework is evaluated through both forward field reconstruction and inverse parameter discovery tasks across a variety of constitutive models, including Oldroyd-B and Giesekus fluids. The results demonstrate the capability of the PINN solver to act as a "virtual rheometer," accurately recovering physical states and parameters in complex, recirculating flows. Finally, this work maps the fundamental limits of the method, characterizing the physical invariances and identifiability boundaries that govern the successful inversion of viscoelastic properties in closed kinematic systems.

Contents

1 Introduction **4**

2 Viscoelastic Fluid Mechanics and the Four-Roll Mill **9**
2.1 General Conservation Laws 9
2.1.1 Mass Balance 10
2.1.2 Momentum Balance 10
2.2 Physical Assumptions 11
2.2.1 Stream Function Formulation 12
2.3 Constitutive Modeling 13
2.3.1 Upper-Convected Time Derivative 13
2.3.2 Oldroyd-B Constitutive Model 14
2.3.3 Giesekus Model 15
2.3.4 Linear Phan-Thien-Tanner (PTT) Model 16
2.4 Characteristic Scales and Dimensionless Numbers 17
2.4.1 Dimensionless Governing System 19
2.5 Four-Roll Mill 19
2.5.1 Flow Index Parameter 20
2.5.2 Stagnation Point 22
2.5.3 Geometric Setup 22
2.5.4 Operating Parameters 22
2.5.5 Boundary Conditions 24

3 Physics-Informed Neural Networks in Complex Rheology **26**
3.1 Fundamentals of Deep Learning 27
3.1.1 The Artificial Neuron and Non-Linear Activation Functions 27
3.1.2 Multi-Layer Perceptron Architecture 28
3.1.3 Loss Functions and Optimization Algorithms 30
3.1.4 Backpropagation 33
3.2 Neural Networks: Capabilities and Physical Limitations 33
3.2.1 Key Advantages 33
3.2.2 Limits of Purely Data-Driven Machine Learning 34
3.3 Physics-Informed Neural Network 35
3.3.1 Physics-Constrained Loss Formulation 35
3.3.2 PINN-Specific Features 38
3.4 Foundational Literature Benchmarks 39

3.4.1 Navier–Stokes Flow Past a Cylinder . . . 40
3.4.2 ViscoelasticNet Framework . . . 41

4 Numerical Results and Methodology 43
4.1 Problem Formulation . . . 44
4.1.1 Recirculating Cavity Challenges . . . 44
4.1.2 Direct and Inverse Modes . . . 44
4.2 Reference CFD Dataset . . . 45
4.2.1 Finite Element Setup . . . 45
4.2.2 Mesh Convergence Study . . . 46
4.3 Network Architecture and Loss Formulation . . . 50
4.3.1 Decoupled Three-Head Network . . . 50
4.3.2 Weight Initialization . . . 51
4.3.3 Governing Residuals and Scaling . . . 51
4.3.4 Loss Formulation . . . 53
4.3.5 Roller Stress Supervision . . . 55
4.4 Training Strategy . . . 56
4.4.1 Two-Stage Physical Decoupling . . . 56
4.4.2 Dual-Precision Optimization . . . 57
4.4.3 Transfer Learning . . . 57
4.5 Direct Results . . . 58
4.5.1 Velocity and Stress . . . 59
4.5.2 Pressure Recovery . . . 60
4.6 Inverse Parameter Discovery . . . 64
4.6.1 Oldroyd-B and Giesekus Benchmarks . . . 65
4.6.2 Transfer Learning Performance . . . 69
4.6.3 Collocation Density Sensitivity . . . 70
4.7 Physical Limits and Identifiability Boundaries . . . 72
4.7.1 Creeping Flow Hydrodynamic Invariance . . . 72
4.7.2 Identifiability of Linear PTT . . . 73

5 Conclusions and Future Perspectives 75

Chapter 1

Introduction

Viscoelastic fluids are present across modern industry and everyday consumer life. Household products, such as toothpaste, mayonnaise, cosmetic gels, and pharmaceutical creams, alongside industrial polymer solutions, melts, and biological suspensions, exhibit rheological behaviors fundamentally more complex than classical Newtonian liquids such as water or light oils [1, 2]. While the deviatoric stress in water responds instantaneously and linearly to the local rate of deformation, viscoelastic materials display elasticity and a fading memory of their deformation history. Under flow, internal macromolecular chains and microstructure stretch, align, and store elastic energy, relaxing over finite timescales toward their equilibrium conformations. Consequently, the local state of stress depends not only on the instantaneous velocity gradients, but integrates the entire deformation trajectory experienced by fluid particles.

Describing the combined effects of viscous dissipation and elastic memory requires governing equations that involve non-trivial, non-linear mathematics [3]. The hyperbolic transport of stress tensors, the necessity of material frame-indifferent rates such as the Upper-Convected Time Derivative (UCTD), and the complex stress saturation closures of non-linear constitutive models present a fascinating field of study. In industrial operations, such as polymer extrusion, in-

jection molding, film coating, and dispersive micro-mixing, these fluids undergo intense, coupled shearing and extensional deformations. In these regimes, flow-induced elastic stresses can trigger macroscopic instabilities, including die swell, melt fracture, and elastic turbulence, which directly govern product throughput and material quality.

Characterizing viscoelastic materials is difficult because standard laboratory rheometers differ kinematically from industrial flow geometries. Conventional rotational rheometers test fluids almost exclusively in simple shear, where polymer chains rotate without significant stretching. As a result, shear data alone cannot capture the extensional response that dominates processing flows [1]. A comprehensive study of viscoelastic fluids requires experimental and numerical geometries that sustain controlled extensional kinematics away from solid walls while allowing mixed flow regimes.

To address this challenge, we introduce the four-roll mill geometry, originally presented by Taylor [4] and formalized by Fuller, Leal, and Bentley [5–7]. The apparatus consists of a closed square cavity containing four cylindrical rollers disposed symmetrically around the center. Each roller can rotate independently, generating a wide variety of planar two-dimensional flow fields inside the cavity, particularly within the central stagnation region. Denoting the angular velocity of diagonal rollers 1 and 3 as Ω and rollers 2 and 4 as $-k\Omega$, adjusting the kinematic parameter k between -1 and $+1$ generates a range of flow kinematics: pure rotation ($k = -1$), simple shear ($k = 0$) with equal strain rate and vorticity, and pure planar extension ($k = +1$) where counter-rotating cylinders create a central hyperbolic stagnation point with pure elongation along the outflow axis. Intermediate values ($-1 < k < 1$) produce mixed shear and extensional fields. This kinematic flexibility makes the four-roll mill a suitable benchmark to study macromolecular deformation and stress accumulation away from wall boundaries.

Machine learning and neural networks have found growing use in computational mechanics for data assimilation and surrogate modeling. Standard data-driven architectures, however, face clear limitations in fluid dynamics. They demand dense training datasets, can produce unphysical predictions that violate mass conservation ($\nabla \cdot \boldsymbol{u} = 0$) or momentum balance, and often fail to generalize outside the sampled domain.

These limitations motivate the adoption of the Physics-Informed Neural Network (PINN) framework, formalized by Raissi, Perdikaris, and Karniadakis [8,9]. Instead of relying on empirical data fitting alone, PINNs embed conservation principles and constitutive differential equations directly into the loss objective through automatic differentiation [10]. Minimizing this composite loss anchors the physical consistency of the predicted fields, reduces the required training data, and improves extrapolation beyond the measurement points.

The necessity of a physics-informed approach becomes particularly evident when addressing inverse rheological problems: determining unknown fluid transport properties and constitutive parameters ($\lambda, \mu_p, \alpha, \varepsilon$) directly from experimental kinematic observations, such as planar velocity fields $\boldsymbol{u}(x, y)$ measured via Particle Image Velocimetry (PIV). Standard Computational Fluid Dynamics (CFD), based on mesh discretization schemes such as the Finite Element Method (FEM) or Finite Volume Method (FVM) [3], is not natively structured for parameter discovery. To identify the physical parameters that best match experimental velocity data, conventional numerical solvers require an external, iterative trial-and-error procedure. Physics-Informed Neural Networks solve this problem directly: the governing differential equations and unknown material parameters are optimized simultaneously during training. Automatic differentiation evaluates the exact sensitivity of the physical residuals with respect to the parameters ($\nabla_{\boldsymbol{\lambda}} \mathcal{L}$), eliminating the need for repeated forward CFD solves and providing a unified

framework for parameter discovery.

Two foundational papers have shaped this research. Raissi et al. [8] first demonstrated the efficacy of PINNs by solving unsteady laminar Newtonian flow past a circular cylinder, demonstrating latent pressure recovery and parameter estimation from velocity data. Then, Thakur, Raissi, and Ardekani [11] extended the methodology to polymeric fluids through the ViscoelasticNet framework, showing that constitutive equations could be coupled with momentum conservation to infer polymeric stresses and extract rheological parameters as a virtual rheometer.

However, ViscoelasticNet was evaluated exclusively on open channel geometries, specifically a stenosed channel and a cross-slot flow. The objective of this dissertation is to extend the physics-informed framework to a closed, recirculating four-roll mill cavity with zero net mass throughput. In this geometry, polymeric stresses are advected along closed streamlines through alternating extensional and shearing zones without convective transport at open boundaries. Furthermore, this work expands the parameter space explored in the literature, testing the framework across higher relaxation times, elevated viscosities, and strong nonlinear mobility regimes up to the theoretical limit $\alpha = 0.50$, while analyzing the physical limits of inverse identifiability.

To accomplish this objective, this thesis develops, validates, and analyzes a physics-informed neural network framework engineered for continuous latent field reconstruction and inverse rheological characterization in the four-roll mill. The architecture comprises three specialized, decoupled Multi-Layer Perceptrons representing the stream function ψ, hydrodynamic pressure p, and polymeric extra-stress tensor $\boldsymbol{\tau}$. Parameterizing the velocity field through $\psi(x, y)$ guarantees exact incompressibility ($\nabla \cdot \boldsymbol{u} \equiv 0$) by construction, while hydrostatic pressure indeterminacy is resolved using an algebraic single-point anchor directly within the

forward pass. Optimization follows a staged schedule, combining initial exploration and second-order quasi-Newton refinement, supported by transfer learning across elasticity regimes.

We validate the solver against finite element solutions generated in COMSOL Multiphysics [12] in two operating modes. In direct mode, the network reconstructs latent stress and pressure fields from planar velocity observations for known rheological parameters. In inverse mode, serving as a virtual rheometer, the model estimates unknown constitutive parameters while reconstructing internal stress fields, tested across different fluid models, collocation densities, and with transfer learning protocol.

Chapter 2

Viscoelastic Fluid Mechanics and the Four-Roll Mill

The hydrodynamics of dilute and semi-dilute polymer solutions fundamentally departs from classical Newtonian fluid behavior. Macromolecular coils dissolved in a solvent undergo continuous conformational changes under applied flow fields, stretching along extensional trajectories and storing elastic energy with a characteristic relaxation memory. Consequently, the local stress state in polymeric fluids integrates the kinematic history experienced by fluid elements, rather than depending only on the instantaneous rate of deformation. Accurately modeling such transport phenomena requires a continuum mechanics formulation where universal conservation laws are coupled with non-linear, frame-indifferent constitutive equations.

2.1 General Conservation Laws

From a macroscopic continuum perspective, fluid motion is governed by the universal conservation laws of mass and linear momentum [1].

2.1.1 Mass Balance

For a fluid possessing local density $\rho(\boldsymbol{x}, t)$, the conservation of mass over an arbitrary control volume is expressed by the continuity equation:

$$\frac{\partial \rho}{\partial t} + \nabla \cdot (\rho \boldsymbol{u}) = 0. \tag{2.1}$$

2.1.2 Momentum Balance

Applying Newton's second law of motion to an infinitesimal fluid element yields the local momentum balance:

$$\rho \frac{\mathrm{D}\boldsymbol{u}}{\mathrm{D}t} = \nabla \cdot \boldsymbol{\sigma} + \rho \boldsymbol{g}, \tag{2.2}$$

where $\boldsymbol{g}$ represents external body force accelerations and $\boldsymbol{\sigma}$ denotes the symmetric total stress tensor. The left-hand side operator $\frac{\mathrm{D}}{\mathrm{D}t}$ is the substantial (or material) derivative, defined as:

$$\frac{\mathrm{D}\boldsymbol{u}}{\mathrm{D}t} \equiv \frac{\partial \boldsymbol{u}}{\partial t} + (\boldsymbol{u} \cdot \nabla)\boldsymbol{u}. \tag{2.3}$$

In the Eulerian framework, the substantial derivative expresses the total rate of velocity change experienced by a moving fluid parcel, including the time derivative $\frac{\partial \boldsymbol{u}}{\partial t}$ at a fixed spatial position and the convective part $(\boldsymbol{u} \cdot \nabla)\boldsymbol{u}$. Expanding Equation (2.2):

$$\rho \left(\frac{\partial \boldsymbol{u}}{\partial t} + \boldsymbol{u} \cdot \nabla \boldsymbol{u} \right) = \nabla \cdot \boldsymbol{\sigma} + \rho \boldsymbol{g}. \tag{2.4}$$

In order to separate isotropic hydrostatic effects from deviatoric shearing and stretching deformations, the total stress tensor $\boldsymbol{\sigma}$ is split into an isotropic pressure contribution and an extra-stress tensor $\boldsymbol{T}$:

$$\boldsymbol{\sigma} = -p\boldsymbol{I} + \boldsymbol{T}, \tag{2.5}$$

where p is the isotropic hydrodynamic pressure, $\boldsymbol{I}$ is the identity tensor, and $\boldsymbol{T}$ represents all viscous and viscoelastic stresses. Substituting Equation (2.5) into Equation (2.4):

$$\rho\left(\frac{\partial \boldsymbol{u}}{\partial t} + \boldsymbol{u}\cdot\nabla\boldsymbol{u}\right) = -\nabla p + \nabla\cdot\boldsymbol{T} + \rho\boldsymbol{g}. \tag{2.6}$$

2.2 Physical Assumptions

The general conservation equations can be simplified by introducing the scale-appropriate physical hypotheses of the four-roll mill:

1. **Incompressibility**: Under moderate operating pressures and low characteristic velocities, density variations are negligible ($\rho = \text{const}$). The continuity equation (2.1) simplifies to the divergence-free condition:

$$\nabla\cdot\boldsymbol{u} = \frac{\partial u}{\partial x} + \frac{\partial v}{\partial y} = 0. \tag{2.7}$$

2. **Steady-state regime**: We focus on the fully established operating regime, setting all explicit time derivatives to zero:

$$\frac{\partial(\cdot)}{\partial t} = \boldsymbol{0}. \tag{2.8}$$

3. **Absence of body forces**: The four-roll mill operates in a horizontal cross-sectional plane (x, y). Because gravity acts along the vertical cylinder axis (z), perpendicular to the flow plane, external body forces vanish: $\boldsymbol{g} = \boldsymbol{0}$.

4. **Isothermal operating conditions**: Because velocities are low and viscous heating is negligible, the temperature field remains uniform. Fluid transport properties, including viscosities and relaxation times, are therefore treated as constants.

5. **Two-dimensional planar kinematics**: Given the large aspect ratio of the cylindrical rollers along the z-axis, three-dimensional end effects are negligible in the transverse mid-plane. The velocity field is modeled as strictly planar:

$$\boldsymbol{u} = \begin{pmatrix} u(x,y) \\ v(x,y) \\ 0 \end{pmatrix}. \tag{2.9}$$

2.2.1 Stream Function Formulation

In two-dimensional incompressible flow, the solenoidal constraint $\nabla \cdot \boldsymbol{u} = 0$ can be satisfied directly by defining a scalar stream function $\psi(x,y)$ such that:

$$u = \frac{\partial \psi}{\partial y}, \qquad v = -\frac{\partial \psi}{\partial x}. \tag{2.10}$$

Because mixed partial derivatives commute ($\frac{\partial^2 \psi}{\partial x \partial y} = \frac{\partial^2 \psi}{\partial y \partial x}$), the velocity divergence vanishes identically:

$$\nabla \cdot \boldsymbol{u} = \frac{\partial u}{\partial x} + \frac{\partial v}{\partial y} = \frac{\partial^2 \psi}{\partial x \partial y} - \frac{\partial^2 \psi}{\partial y \partial x} \equiv 0. \tag{2.11}$$

Parameterizing the velocity through ψ satisfies mass conservation everywhere by construction, replacing the two velocity components (u, v) with a single scalar field.

Under these physical assumptions and kinematically reduced velocity representation, the steady balance of linear momentum simplifies to:

$$\rho(\boldsymbol{u} \cdot \nabla \boldsymbol{u}) = -\nabla p + \nabla \cdot \boldsymbol{T}. \tag{2.12}$$

2.3 Constitutive Modeling

In two dimensions, the momentum balance (2.12) contains six unknown scalar fields $(u, v, p, T_{xx}, T_{xy}, T_{yy})$, while mass and momentum conservation provide only three equations. A constitutive equation for the extra-stress tensor $\boldsymbol{T}$ is required to close the system.

For dilute and semi-dilute polymer solutions, we split the extra-stress tensor $\boldsymbol{T}$ into a Newtonian solvent contribution and an elastic polymeric contribution $\boldsymbol{\tau}$ [1, 2]:

$$\boldsymbol{T} = 2\mu_s \boldsymbol{D} + \boldsymbol{\tau}, \tag{2.13}$$

where μ_s is the dynamic viscosity of the Newtonian solvent, and $\boldsymbol{D}$ is the symmetric rate-of-strain tensor:

$$\boldsymbol{D} = \frac{1}{2}\left[\nabla \boldsymbol{u} + (\nabla \boldsymbol{u})^T\right]. \tag{2.14}$$

Taking the divergence of the solvent stress and leveraging incompressibility ($\nabla \cdot \boldsymbol{u} = 0$):

$$\nabla \cdot (2\mu_s \boldsymbol{D}) = \mu_s \nabla \cdot \left(\nabla \boldsymbol{u} + (\nabla \boldsymbol{u})^T\right) = \mu_s \left(\nabla^2 \boldsymbol{u} + \nabla(\nabla \cdot \boldsymbol{u})\right) = \mu_s \nabla^2 \boldsymbol{u}. \tag{2.15}$$

Substituting Equation (2.15) into Equation (2.12) yields the steady momentum equation:

$$\rho(\boldsymbol{u} \cdot \nabla \boldsymbol{u}) = -\nabla p + \mu_s \nabla^2 \boldsymbol{u} + \nabla \cdot \boldsymbol{\tau}. \tag{2.16}$$

2.3.1 Upper-Convected Time Derivative

Polymer chains suspended in a solvent stretch, orient, and relax over finite timescales. Because fluid parcels translate, rotate, and deform as they move through the domain, constitutive relations require a rate of stress that is invariant under changes

of reference frame (material frame-indifference). A standard time derivative evaluated at a fixed point would interpret rigid-body rotation or parcel stretching as a change in stress, even if the internal state of the polymer remained unaltered.

To measure true rheological evolution, the time rate of change must be evaluated in a local coordinate system that deforms affinely with the fluid element. This objective tensor rate is the *Upper-Convected Time Derivative* (UCTD), denoted by $\overset{\nabla}{\boldsymbol{\tau}}$:

$$\overset{\nabla}{\boldsymbol{\tau}} \equiv \frac{\partial \boldsymbol{\tau}}{\partial t} + \boldsymbol{u} \cdot \nabla \boldsymbol{\tau} - (\nabla \boldsymbol{u})\boldsymbol{\tau} - \boldsymbol{\tau}(\nabla \boldsymbol{u})^T, \tag{2.17}$$

where $(\nabla \boldsymbol{u})_{ij} = \partial u_i / \partial x_j$. For steady flows ($\partial \boldsymbol{\tau}/\partial t = \mathbf{0}$), the operator combines advective stress transport along streamlines, $\boldsymbol{u} \cdot \nabla \boldsymbol{\tau}$, with kinematic stretching and rotation induced by local velocity gradients, $-(\nabla \boldsymbol{u})\boldsymbol{\tau} - \boldsymbol{\tau}(\nabla \boldsymbol{u})^T$.

2.3.2 Oldroyd-B Constitutive Model

The Oldroyd-B model [13] provides the standard starting point for viscoelastic fluid modeling. It idealizes macromolecules as non-interacting Hookean dumbbells suspended in a Newtonian solvent [1]. Under steady conditions, the polymeric extra-stress tensor $\boldsymbol{\tau}$ satisfies:

$$\boldsymbol{\tau} + \lambda \overset{\nabla}{\boldsymbol{\tau}} = 2\mu_p \boldsymbol{D}, \tag{2.18}$$

where λ is the relaxation time and μ_p is the polymeric viscosity. The total zero-shear dynamic viscosity of the solution is:

$$\mu_0 = \mu_s + \mu_p. \tag{2.19}$$

In steady simple shear flow, Oldroyd-B predicts a constant viscosity $\eta(\dot{\gamma}) = \mu_0$, reproducing the behavior of Boger fluids.

Planar Extensional Behavior

In strong extensional flows, the Oldroyd-B model exhibits an unphysical stress singularity. For a steady planar elongational flow centered at the origin, $u = Ex$ and $v = -Ey$ with extension rate $E > 0$, the tensile normal stress along the stretching axis is:

$$\tau_{xx} = \frac{2\mu_p E}{1 - 2\lambda E}. \tag{2.20}$$

As the extension rate approaches the critical threshold:

$$E_{\text{crit}} = \frac{1}{2\lambda}, \tag{2.21}$$

the denominator vanishes. Because Hookean springs have infinite extensibility, the polymer chains can stretch without bound, driving both τ_{xx} and the extensional viscosity η_E to infinity.

This stress singularity near $2\lambda E \to 1$ is the constitutive origin of the *High-Weissenberg Number Problem* (HWNP). Near extensional stagnation points, it generates steep stress layers that cause numerical instability in standard mesh-based solvers [3].

2.3.3 Giesekus Model

To overcome the unphysical stress singularity of Oldroyd-B, Giesekus [1, 14] introduced a quadratic non-linear stress contribution derived from configuration-dependent hydrodynamic drag. In semi-dilute solutions, stretching macromolecular chains align parallel to each other, reducing drag along the orientation direction relative to the transverse directions. This anisotropic friction is incorporated by modifying the stress relaxation rate:

$$\boldsymbol{\tau} + \lambda \overset{\nabla}{\boldsymbol{\tau}} + \frac{\alpha\lambda}{\mu_p}\boldsymbol{\tau}^2 = 2\mu_p \boldsymbol{D}, \tag{2.22}$$

where $\alpha \in [0, 0.5]$ is the dimensionless mobility parameter, and $\boldsymbol{\tau}^2 = \boldsymbol{\tau} \cdot \boldsymbol{\tau}$. When $\alpha = 0$, isotropic drag is recovered and the model reduces identically to Oldroyd-B. For $\alpha > 0$, the quadratic dissipation term $\boldsymbol{\tau}^2$ provides two key rheological enhancements:

1. **Bounded extensional plateau**: Extensional stresses no longer diverge at E_{crit}; instead, η_E reaches an asymptotic finite plateau at high strain rates, completely resolving the constitutive singularity.

2. **Shear-thinning behavior**: The steady shear viscosity $\eta(\dot{\gamma})$ decreases monotonically with increasing shear rate, closely matching experimental polymer melt and concentrated solution data.

2.3.4 Linear Phan-Thien-Tanner (PTT) Model

Derived from the transient network theory of concentrated polymer solutions and melts [1,15], the Phan-Thien-Tanner (PTT) model treats macromolecules as a network of temporary entanglements. Under strong deformation, mechanical tension on the network strands accelerates junction destruction. In its linear formulation, this stress-induced network breakdown is modeled by scaling the linear relaxation term:

$$f(\boldsymbol{\tau})\boldsymbol{\tau} + \lambda \overset{\nabla}{\boldsymbol{\tau}} = 2\mu_p \boldsymbol{D}, \tag{2.23}$$

where the scalar function $f(\boldsymbol{\tau})$ is defined as:

$$f(\boldsymbol{\tau}) = 1 + \frac{\varepsilon\lambda}{\mu_p}\text{tr}(\boldsymbol{\tau}), \tag{2.24}$$

with $\text{tr}(\boldsymbol{\tau}) = \tau_{xx} + \tau_{yy}$ (since $\tau_{zz} = 0$ under planar deformation for linear PTT) and $\varepsilon \geq 0$ representing the dimensionless extensibility parameter. When $\varepsilon = 0$, $f(\boldsymbol{\tau}) \equiv 1$ and the linear Oldroyd-B constitutive law is recovered. For $\varepsilon > 0$, high

trace stresses accelerate stress relaxation, bounding extensional viscosity growth and preventing unbounded coil stretch near extensional stagnation points.

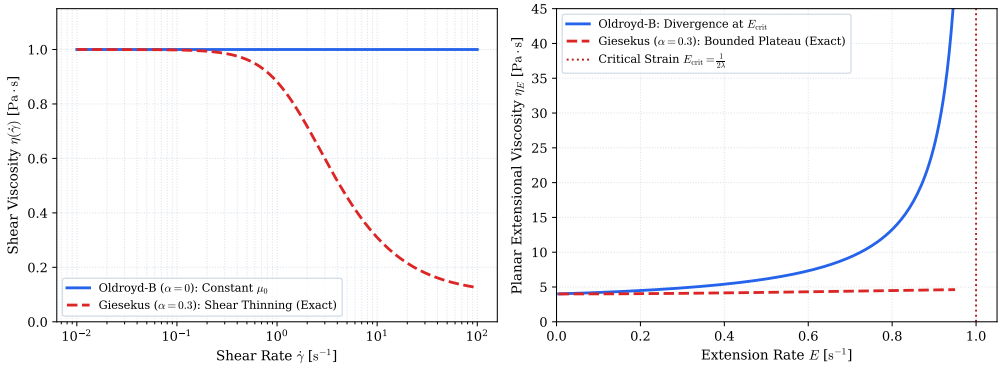


Figure 2.1: Example comparison of rheological model responses: (a) steady shear viscosity $\eta(\dot{\gamma})$ as a function of shear rate, illustrating the constant Newtonian-like plateau of Oldroyd-B ($\alpha = 0$) versus the non-linear shear-thinning response of Giesekus ($\alpha > 0$); (b) planar extensional viscosity $\eta_E(E)$ as a function of extension rate, contrasting the unphysical asymptotic divergence of Oldroyd-B at the critical strain threshold $E_{\text{crit}} = 1/(2\lambda)$ with the bounded saturation plateau provided by non-linear closures.

2.4 Characteristic Scales and Dimensionless Numbers

To characterize the flow regimes in the four-roll mill, we define characteristic kinematic and dynamic scales alongside the governing dimensionless groups of viscoelastic fluid mechanics [1, 3].

We select the roller radius as the characteristic transverse geometric scale, $H_{\text{ref}} = R$, the roller surface peripheral speed as the characteristic velocity scale, $U_{\text{ref}} = u_{\text{roll}}$, and the total zero-shear dynamic viscosity of the solution, $\mu_0 = \mu_s + \mu_p$, as the characteristic viscous scale. The balance of inertial, elastic, and

viscous forces is governed by three fundamental dimensionless groups:

$$Re = \frac{\rho U_{\text{ref}} H_{\text{ref}}}{\mu_0} \qquad \text{(Reynolds number)}, \tag{2.25}$$

$$Wi = \frac{\lambda U_{\text{ref}}}{H_{\text{ref}}} \qquad \text{(Weissenberg number)}, \tag{2.26}$$

$$\beta = \frac{\mu_s}{\mu_0} = \frac{\mu_s}{\mu_s + \mu_p} \qquad \text{(Solvent viscosity ratio)}. \tag{2.27}$$

Each dimensionless group carries a precise physical significance:

- **Reynolds number (Re)**: Compares convective inertial forces to viscous dissipative forces. In the four-roll mill configurations studied here, $Re = 0.0417 \ll 1$, placing the flow within the laminar creeping flow (Stokes) regime where fluid inertia is negligible.

- **Weissenberg number (Wi)**: Compares the characteristic elastic relaxation time of the polymer λ to the kinematic timescale of the flow, $t_{\text{flow}} = H_{\text{ref}}/U_{\text{ref}}$. It quantifies the degree of flow-induced molecular orientation and elastic non-linearity.

- **Solvent viscosity ratio (β)**: Measures the fraction of total solution viscosity contributed by the Newtonian solvent. The limits $\beta = 1$ and $\beta \to 0$ correspond to a purely Newtonian fluid and a solvent-free polymer melt, respectively.

2.4.1 Dimensionless Governing System

Expressed in terms of dimensionless variables, the steady planar governing equations become:

$$\nabla \cdot \boldsymbol{u} = 0, \tag{2.28}$$

$$Re(\boldsymbol{u} \cdot \nabla \boldsymbol{u}) + \nabla p - \beta \nabla^2 \boldsymbol{u} - \nabla \cdot \boldsymbol{\tau} = \boldsymbol{0}, \tag{2.29}$$

$$\boldsymbol{\mathcal{G}}(\boldsymbol{\tau}) + Wi\left(\boldsymbol{u} \cdot \nabla \boldsymbol{\tau} - (\nabla \boldsymbol{u})\boldsymbol{\tau} - \boldsymbol{\tau}(\nabla \boldsymbol{u})^T\right) - 2(1-\beta)\boldsymbol{D} = \boldsymbol{0}, \tag{2.30}$$

where the constitutive operator $\boldsymbol{\mathcal{G}}(\boldsymbol{\tau})$ provides a unified closure for the models under investigation:

$$\boldsymbol{\mathcal{G}}(\boldsymbol{\tau}) = \begin{cases} \boldsymbol{\tau} & \text{(Oldroyd-B)}, \\ \boldsymbol{\tau} + \frac{\alpha Wi}{1-\beta}\boldsymbol{\tau}^2 & \text{(Giesekus)}, \\ \left(1 + \frac{\varepsilon Wi}{1-\beta}\mathrm{tr}(\boldsymbol{\tau})\right)\boldsymbol{\tau} & \text{(Linear PTT)}. \end{cases} \tag{2.31}$$

2.5 Four-Roll Mill

The four-roll mill is a classic benchmark designed to generate controlled two-dimensional planar extensional flow with a central stagnation point [4–7]. Originally introduced by Taylor [4] and systematically analyzed by Fuller and Leal [5,6], it provides an ideal apparatus for investigating macromolecular deformation away from solid walls. As illustrated conceptually in Figure 2.2, the setup comprises four cylindrical rollers placed symmetrically inside a closed square cavity.

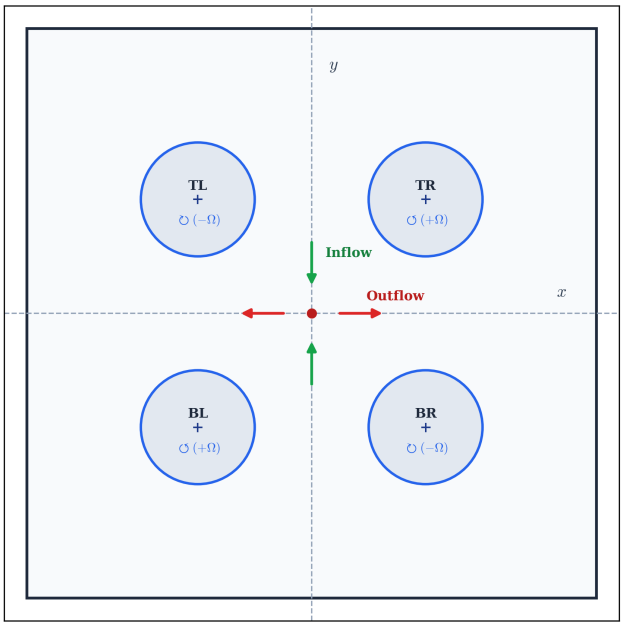


Figure 2.2: Conceptual diagram of the four-roll mill geometry, illustrating the four rotating cylinders (adjacent rolls counter-rotate with angular velocity Ω, diagonal rolls co-rotate), the central hyperbolic stagnation point $(0, 0)$, and the resulting elongational inflow and outflow directions.

2.5.1 Flow Index Parameter

In the central region bounded by the four rollers, the velocity field can be represented as a linear superposition of extension and solid-body rotation [5, 6]:

$$\boldsymbol{u}(x, y) = \begin{pmatrix} Ex + \omega y \\ -Ey - \omega x \end{pmatrix}, \tag{2.32}$$

where $E \geq 0$ is the extension rate and $\omega \geq 0$ is the angular rotation rate. To quantify the relative magnitude of extension versus rotation, we define a dimensionless flow type parameter $\chi \in [-1, 1]$:

$$\chi = \frac{E - \omega}{E + \omega}. \tag{2.33}$$

Depending on roller rotation speeds and directions, this local kinematics covers three canonical flow regimes (in Figure 2.3):

- $\chi = +1$ ($\omega = 0$): **Pure Planar Extensional Flow**. Generated when adjacent rollers counter-rotate symmetrically (diagonal rollers co-rotate). Fluid elements experience pure elongation along x and compression along y, producing hyperbolic streamlines $xy = \text{const}$.

- $\chi = 0$ ($E = \omega$): **Simple Shear Flow**. Characterized by equal strain and rotation rates, producing rectilinear shear streamlines locally near the origin.

- $\chi = -1$ ($E = 0$): **Pure Solid-Body Rotation**. Generated when rollers rotate in unison to drive bulk cavity circulation, resulting in concentric circular streamlines without stretching.

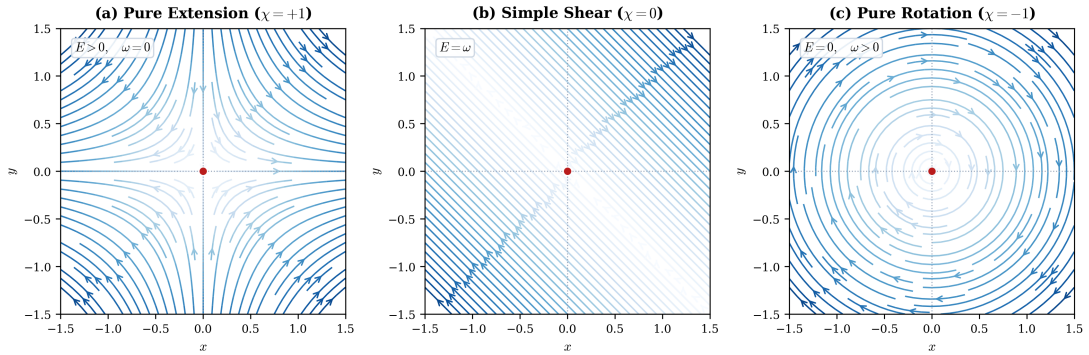


Figure 2.3: Streamline patterns of the two-dimensional linear flow field in Equation (2.32) for the three canonical regimes of the flow classification parameter χ: (a) pure planar extension ($\chi = +1$), (b) simple shear ($\chi = 0$), and (c) pure solid-body rotation ($\chi = -1$). The red marker indicates the central stagnation point $(0, 0)$.

2.5.2 Stagnation Point

In the planar extensional regime ($\chi = +1$), the velocity vanishes identically at the cavity center:

$$\boldsymbol{u}(0,0) = \mathbf{0}. \tag{2.34}$$

Along the vertical inflow streamline ($x = 0$), incoming fluid parcels decelerate toward zero velocity, reaching asymptotically long residence times near the origin. This kinematics allows elastic strain to accumulate over durations comparable to or exceeding the polymer relaxation time λ, making the four-roll mill an ideal benchmark for probing strong extensional stresses far from solid walls.

2.5.3 Geometric Setup

Following the established benchmark setup of D'Avino et al. [16], we define the physical dimensions of the apparatus:

- **Square cavity**: Side length $L = 50\,\text{mm}$ spanning the domain $[-25, 25]\,\text{mm}$.
- **Cylindrical rollers**: Radius $R = 5\,\text{mm}$ ($H_{\text{ref}} = R = 0.1L$).
- **Roller centers**: Located symmetrically at $(\pm 10, \pm 10)\,\text{mm}$ ($d_{\text{roll}} = 10\,\text{mm}$).
- **Dimensionless domain**: In units of $H_{\text{ref}} = R$, the computational domain spans $[-5, 5]^2$, with roller centers located at $(\pm 2.0, \pm 2.0)$ and non-dimensional roller radius $R^* = 1.0$.

As shown in Figure 2.4, these dimensions define the operational space and solid boundaries of the cavity.

2.5.4 Operating Parameters

The flow is driven by the rotation of the four cylindrical rollers with peripheral surface velocity $u_{\text{roll}} = \Omega R = 8.333 \times 10^{-3}\,\text{m/s}$. The geometric, fluid physical,

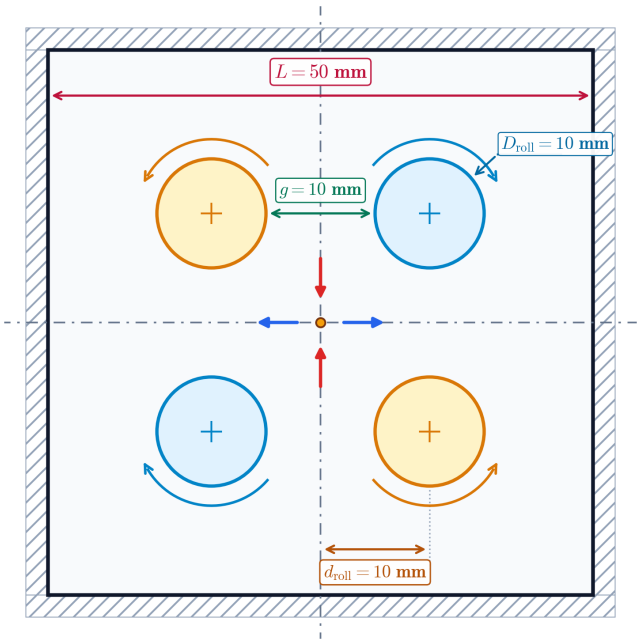


Figure 2.4: Detailed geometric dimensions and boundary partition of the four-roll mill: planar fluid domain Ω, external stationary cavity boundary Γ_{outer} with side length $L = 50\,\text{mm}$, four cylinder boundaries $\Gamma_{\text{rolls}} = \bigcup_k \Gamma_k$ ($k \in \{\text{TL}, \text{TR}, \text{BL}, \text{BR}\}$) with radius $R = 5\,\text{mm}$ and center offset $d_{\text{roll}} = 10\,\text{mm}$, peripheral rotation directions with angular velocity Ω, and central extensional stagnation point $(0, 0)$.

and characteristic dimensionless parameters investigated across the benchmark test cases are organized systematically in Table 2.1.

Table 2.1: Geometric, physical, and dimensionless parameters of the four-roll mill benchmark.

Parameter	**Symbol**	**Value / Range**	**Unit**
Geometric and Kinematic Setup			
Cavity side length	L	0.05	m
Roller radius (reference length)	$H_{\text{ref}} = R$	0.005	m
Roller center offset	d_{roll}	0.010	m
Roller peripheral velocity	$u_{\text{roll}} = U_{\text{ref}}$	8.333×10^{-3}	m/s
Fluid Physical Properties			
Fluid density	ρ	1000	kg/m^3
Characteristic viscosity scale	$\eta_0 = \mu_0$	1.0	$\text{Pa} \cdot \text{s}$
Solvent dynamic viscosity	μ_s	0.1–0.5	$\text{Pa} \cdot \text{s}$
Polymer dynamic viscosity	μ_p	0.5–0.9	$\text{Pa} \cdot \text{s}$
Polymer relaxation time	λ	0.05–1.0	s
Dimensionless Governing Groups			
Reynolds number	Re	$0.0417 \ll 1$	–
Weissenberg number	Wi	0.083–1.667	–
Viscosity ratio	β	0.1–0.5	–

2.5.5 Boundary Conditions

Let Ω denote the planar fluid domain bounded externally by the square cavity walls Γ_{outer} and internally by the boundaries of the four cylindrical rollers $\Gamma_{\text{rolls}} = \bigcup_k \Gamma_k$ ($k \in \{\text{TL}, \text{TR}, \text{BL}, \text{BR}\}$), such that the total boundary is given by the union $\partial\Omega = \Gamma_{\text{outer}} \cup \Gamma_{\text{rolls}}$, as illustrated in Figure 2.4.

To close the mathematical problem formulated by mass conservation (2.7), momentum balance (2.16), and the constitutive equations, kinematic Dirichlet conditions and a pressure reference point are prescribed:

$$\boldsymbol{u}(\boldsymbol{x}) = \boldsymbol{0}, \qquad \boldsymbol{x} \in \Gamma_{\text{outer}}, \tag{2.35}$$

$$\boldsymbol{u}(\boldsymbol{x}) = \boldsymbol{\Omega}_k \times (\boldsymbol{x} - \boldsymbol{x}_k), \qquad \boldsymbol{x} \in \Gamma_k \quad (k \in \{\text{TL}, \text{TR}, \text{BL}, \text{BR}\}), \tag{2.36}$$

$$p(\boldsymbol{x}_0) = 0, \qquad \boldsymbol{x}_0 \in \Omega. \tag{2.37}$$

In the planar extensional regime ($\chi = +1$), the roller angular velocity vectors $\boldsymbol{\Omega}_k = (0, 0, \Omega_k)^T$ are oriented along the out-of-plane z-axis such that diagonal pairs co-rotate and adjacent pairs counter-rotate:

$$\Omega_{\mathrm{TL}} = -\Omega, \quad \Omega_{\mathrm{TR}} = +\Omega, \quad \Omega_{\mathrm{BL}} = +\Omega, \quad \Omega_{\mathrm{BR}} = -\Omega, \tag{2.38}$$

with peripheral angular speed $\Omega = u_{\mathrm{roll}}/R$.

Because fluid cannot penetrate the impermeable solid boundaries ($\boldsymbol{u} \cdot \boldsymbol{n} = 0$ on $\partial\Omega$), the global mass conservation constraint $\oint_{\partial\Omega} \boldsymbol{u} \cdot \boldsymbol{n} \, \mathrm{d}\Gamma = 0$ is satisfied identically. The condition (2.37) anchors the hydrostatic pressure field. Finally, because the four-roll mill constitutes a completely closed recirculating cavity without inflow boundaries ($\Gamma_{\mathrm{in}} = \emptyset$), no external Dirichlet conditions are mathematically required for the extra-stress tensor $\boldsymbol{\tau}$.

Chapter 3

Physics-Informed Neural Networks in Complex Rheology

Feedforward deep neural networks, commonly known as Multi-Layer Perceptrons (MLPs), provide flexible continuous approximations for non-linear functions [17]. An MLP is formed by composing a sequence of parametric functions, called *layers*. Each layer maps its inputs through a linear transformation followed by a non-linear operation. With these layers in series, the network can represent continuous functional mappings of arbitrary complexity without requiring explicit basis functions.

In scientific computing, neural networks represent physical fields directly as continuous functions over spatial coordinates, rather than storing discrete state variables on a mesh. However, using purely data-driven networks for fluid mechanics presents clear limits: they require large amounts of labeled data, can violate conservation laws, and cannot infer latent material properties on their own.

Physics-Informed Neural Networks (PINNs) address these limits by embedding governing partial differential equations directly into the training loss. The differential system formulated in Chapter 2, governed by mass conservation, vis-

coelastic momentum balance, and frame-indifferent constitutive transport, represents a coupled set of non-linear equations that poses well-known challenges to traditional mesh-based methods. In this chapter, we formulate this continuous boundary-value problem within a Scientific Machine Learning (SciML) framework. By embedding momentum and constitutive balances into the optimization objective through automatic differentiation, the PINN framework enables both direct state reconstruction and inverse rheological parameter discovery without spatial discretization meshes.

3.1 Fundamentals of Deep Learning

Deep learning models approximate continuous mathematical functions by learning parameters directly from data [17]. In this section, we review the fundamental building blocks of feedforward networks: the artificial neuron, activation functions, multi-layer architectures, loss functions, optimization algorithms, and gradient computation.

3.1.1 The Artificial Neuron and Non-Linear Activation Functions

The elementary computational unit of a neural network is the artificial neuron. Given an input vector $\boldsymbol{x}$, a neuron computes a linear combination of the inputs followed by a non-linear activation:

$$z = \boldsymbol{w} \cdot \boldsymbol{x} + b, \qquad a = \sigma(z), \tag{3.1}$$

where $\boldsymbol{w}$ is the weight vector, b is a scalar bias, $\sigma(\cdot)$ is the activation function and a is the output. Non-linear activations are essential: without them, composing multiple linear layers simply collapses into a single linear transformation.

In classical machine learning, piecewise-linear functions such as the Rectified Linear Unit, $\mathrm{ReLU}(z) = \max(0, z)$, are widely favored for their computational efficiency. However, in scientific computing and fluid mechanics, the network outputs parameterize continuous physical fields governed by differential equations. These governing balance laws involve spatial derivatives of first, second, or third order, such as viscous diffusion $\mu_s \nabla^2 \boldsymbol{u}$ or stress divergence $\nabla \cdot \boldsymbol{\tau}$.

Because the first derivative of ReLU is discontinuous and its second derivative is zero almost everywhere, evaluating a ReLU network inside differential residuals eliminates diffusive and curvature terms. Physics-informed models therefore require smooth, differentiable activation functions ($\mathcal{C}^2$ or higher). In this work, we employ the Sigmoid Linear Unit (SiLU, also known as Swish) [17]:

$$\mathrm{SiLU}(z) = \frac{z}{1 + e^{-z}}. \tag{3.2}$$

The SiLU activation is smooth and non-monotonic. Unlike the hyperbolic tangent (tanh), which saturates toward constant values for large positive inputs, SiLU remains active in the positive regime. This property maintains healthy gradient propagation in flow regions with intense shear and steep localized stress gradients, which are characteristic of viscoelastic flows.

3.1.2 Multi-Layer Perceptron Architecture

By arranging artificial neurons into sequential layers, we construct a Multi-Layer Perceptron (MLP). The network maps an input vector $\boldsymbol{x}$ to an output prediction

$\hat{\boldsymbol{u}}$ by propagating information through L successive layers [17]:

$$\begin{aligned} \boldsymbol{h}^{(1)} &= \sigma\left(\boldsymbol{W}^{(1)}\boldsymbol{x} + \boldsymbol{b}^{(1)}\right), \\ \boldsymbol{h}^{(l)} &= \sigma\left(\boldsymbol{W}^{(l)}\boldsymbol{h}^{(l-1)} + \boldsymbol{b}^{(l)}\right), \quad \text{for } l = 2, \ldots, L-1, \\ \hat{\boldsymbol{u}} &= \boldsymbol{W}^{(L)}\boldsymbol{h}^{(L-1)} + \boldsymbol{b}^{(L)}, \end{aligned} \tag{3.3}$$

where $\boldsymbol{h}^{(l)}$ denotes the hidden feature vector at layer l, and $\boldsymbol{W}^{(l)}$ and $\boldsymbol{b}^{(l)}$ represent the weight matrix and bias vector of that layer. In our implementation, the inputs $\boldsymbol{x} = (x, y)$ are spatial coordinates within the fluid cavity, while the outputs $\hat{\boldsymbol{u}}$ parameterize the physical state variables: the stream function ψ, hydrodynamic pressure p, and extra-stress tensor components $\boldsymbol{\tau}$.

A central theoretical property of MLPs is universal function approximation [18,19]: a feedforward network with sufficient width and smooth non-linear activations can approximate any continuous function on a compact domain to arbitrary precision. However, universal approximation guarantees only the existence of a suitable set of weights, not that standard gradient-based optimization will automatically find them or that spatial derivatives will match the physical solution with equal fidelity.

Importance of Weight Initialization

The initialization of network parameters strongly conditions the resulting optimization dynamics. Because deep architectures compose multiple transformations and non-linearities, inappropriate initial weight magnitudes can cause layer activations and backpropagated gradients to either vanish or explode across depth [17]. In scientific computing and differential problems, where non-linear convective terms and tensor couplings are evaluated, uncalibrated initial weights can trigger extreme initial residual spikes that immediately destabilize gradient

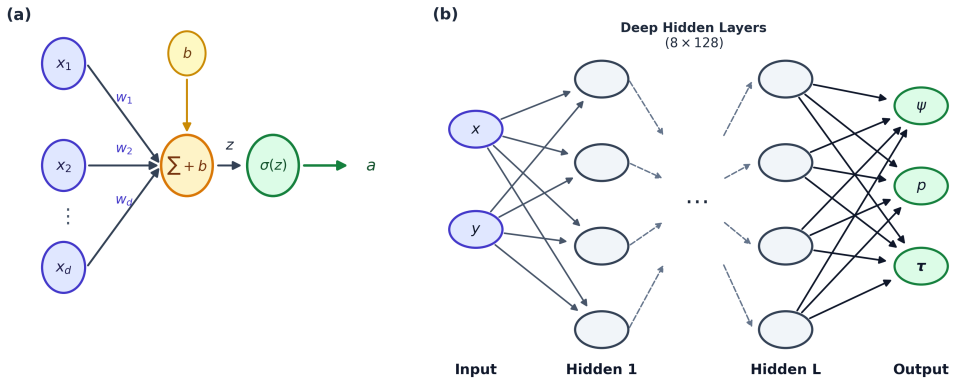


Figure 3.1: Neural network representations: (a) schematic of an artificial neuron computing an affine combination $z = \boldsymbol{w} \cdot \boldsymbol{x} + b$ followed by a smooth activation $\sigma(z)$; (b) deep Multi-Layer Perceptron (MLP) architecture propagating coordinate inputs (x, y) through L hidden layers to parameterize continuous physical fields $(\psi, p, \boldsymbol{\tau})$.

descent. Appropriate variance-preserving initialization schemes, such as scaling parameter distributions inversely with layer widths, help preserve stable gradient flow and avoid optimization divergence. Specific configuration choices regarding hidden and output layer initialization adopted in this work are detailed in Chapter 4.

3.1.3 Loss Functions and Optimization Algorithms

Training an MLP consists of identifying a parameter set $\boldsymbol{\theta} = \{\boldsymbol{W}^{(l)}, \boldsymbol{b}^{(l)}\}$ that minimizes a scalar objective function $\mathcal{L}(\boldsymbol{\theta})$. In supervised regression, the discrepancy between model predictions and target reference measurements is commonly quantified through the Mean Squared Error (MSE):

$$\mathcal{L}_{\text{MSE}}(\boldsymbol{\theta}) = \frac{1}{N} \sum_{i=1}^{N} \left\| \hat{\boldsymbol{u}}(\boldsymbol{x}_i; \boldsymbol{\theta}) - \boldsymbol{u}_i^{\text{target}} \right\|_2^2 . \tag{3.4}$$

Parameters are iteratively updated along descent directions derived from the loss gradient:

$$\boldsymbol{\theta} \leftarrow \boldsymbol{\theta} - \eta \boldsymbol{d}_k, \tag{3.5}$$

where $\eta > 0$ represents the learning rate (step size) and $\boldsymbol{d}_k$ denotes the search direction at iteration k.

In scientific machine learning, gradient-based solvers must navigate non-convex loss surfaces characterized by ill-conditioned curvature and local minima. Two widely adopted algorithms in this setting are:

Adam: First Order Estimation

The Adam optimizer [20] computes individual adaptive learning rates for each network parameter by tracking running estimates of the first and second moments of the loss gradient. At iteration k, given the gradient $\boldsymbol{g}_k = \nabla_{\boldsymbol{\theta}} \mathcal{L}(\boldsymbol{\theta}_k)$, the algorithm updates two vectors recursively:

$$\boldsymbol{m}_k = \beta_1 \boldsymbol{m}_{k-1} + (1 - \beta_1)\boldsymbol{g}_k, \tag{3.6}$$

$$\boldsymbol{v}_k = \beta_2 \boldsymbol{v}_{k-1} + (1 - \beta_2)\boldsymbol{g}_k^2, \tag{3.7}$$

where the square $\boldsymbol{g}_k^2$ is evaluated component-wise, and $\beta_1, \beta_2 \in [0, 1)$ are exponential decay factors (canonically set to $\beta_1 = 0.9$ and $\beta_2 = 0.999$).

These two vectors have distinct physical and optimization roles:

- **First moment vector $\boldsymbol{m}_k$ (Directional Momentum)**: maintains an exponentially decaying average of past gradients. It acts as a directional inertia, accelerating descent along persistent slopes while dampening high-frequency oscillations across walls.

- **Second uncentered moment vector $\boldsymbol{v}_k$ (Adaptive Scaling)**: tracks the running magnitude of squared gradient components. By dividing the update step by $\sqrt{\boldsymbol{v}_k}$, Adam rescales the effective learning rate coordinate by coordinate, taking cautious smaller steps along steep, noisy directions and bolder steps along flat plateaus.

Because both moment vectors are initialized to zero ($\boldsymbol{m}_0 = \boldsymbol{v}_0 = \boldsymbol{0}$), they initially exhibit a numerical bias toward zero. Kingma and Ba [20] rescale $\boldsymbol{m}_k$ and $\boldsymbol{v}_k$ by $(1-\beta_1^k)$ and $(1-\beta_2^k)$ to correct this initial bias before updating the parameter vector $\boldsymbol{\theta}_{k+1}$. This adaptive mechanism makes Adam particularly robust on rugged, multi-objective loss landscapes characterized by disparate gradient scales, serving as an effective global explorer in early training.

L-BFGS: Second-Order Refinement

While first-order methods explore the loss landscape robustly, refining physical solutions and satisfying stiff differential equations requires incorporating second-order curvature information. Classical Newton methods require forming and inverting the full Hessian matrix $\nabla^2_{\boldsymbol{\theta}}\mathcal{L}$, which is computationally prohibitive for neural networks containing $N \sim 10^5$–10^6 parameters ($\mathcal{O}(N^2)$ memory and $\mathcal{O}(N^3)$ inversion cost).

Quasi-Newton methods overcome this limitation by iteratively approximating the inverse Hessian matrix $\boldsymbol{H}_k \approx (\nabla^2_{\boldsymbol{\theta}}\mathcal{L})^{-1}$ using only successive parameter displacements $\boldsymbol{s}_k = \boldsymbol{\theta}_{k+1} - \boldsymbol{\theta}_k$ and gradient variations $\boldsymbol{y}_k = \nabla_{\boldsymbol{\theta}}\mathcal{L}_{k+1} - \nabla_{\boldsymbol{\theta}}\mathcal{L}_k$. The classic BFGS update formula is given by [21]:

$$\boldsymbol{H}_{k+1} = (\boldsymbol{I} - \rho_k \boldsymbol{s}_k \boldsymbol{y}_k^\top)\boldsymbol{H}_k(\boldsymbol{I} - \rho_k \boldsymbol{y}_k \boldsymbol{s}_k^\top) + \rho_k \boldsymbol{s}_k \boldsymbol{s}_k^\top, \qquad \rho_k = \frac{1}{\boldsymbol{y}_k^\top \boldsymbol{s}_k}. \tag{3.8}$$

To avoid storing the dense $N \times N$ matrix $\boldsymbol{H}_{k+1}$, the Limited-memory BFGS (L-BFGS) algorithm [21,22] stores only the m most recent vector pairs $\{\boldsymbol{s}_i, \boldsymbol{y}_i\}_{i=k-m}^{k-1}$ (typically $m \in [10, 50]$). The descent direction is then evaluated on the fly via an efficient two-loop recursion algorithm, dramatically reducing the memory footprint from $\mathcal{O}(N^2)$ to $\mathcal{O}(mN)$. The step length along this direction is calibrated through a line search ensuring sufficient decrease of the objective function. When evaluated over full collocation batches in double precision (FP64), L-BFGS ex-

hibits fast superlinear asymptotic convergence, making it ideal for the final refinement stage of PINN solvers.

3.1.4 Backpropagation

To update network parameters via gradient descent, we compute the partial derivatives of the scalar loss with respect to all weights and biases, $\nabla_{\boldsymbol{\theta}}\mathcal{L}$. Evaluating these gradients through individual numerical perturbations would be computationally prohibitive.

Instead, neural networks compute analytical gradients using the *Automatic Differentiation (AD)* algorithm [10]. Automatic differentiation evaluates derivatives to machine precision by applying the chain rule to elementary operations along the computational graph. In standard machine learning, reverse-mode AD is widely known as the *backpropagation* algorithm, primarily employed to evaluate the gradient of the scalar loss with respect to network parameters.

This reverse-mode automatic differentiation evaluates all parameter gradients simultaneously at a computational cost comparable to a single forward pass, making deep network training feasible.

3.2 Neural Networks: Capabilities and Physical Limitations

3.2.1 Key Advantages

Neural networks offer distinctive characteristics that make them attractive for continuous physical modeling:

1. **High-Dimensional Functional Approximation**: The universal approximation capacity allows neural networks to represent complex, coupled phys-

ical fields (velocities, pressures, and stress tensors) within a unified functional representation, without requiring specialized interpolation elements for each variable.

2. **Meshless Representation**: Traditional grid-based methods (Finite Element, Finite Volume, and Finite Difference Methods) represent state variables exclusively at discrete vertices or cell centers, relying on explicit mesh connectivity. In contrast, a neural network provides a continuous, globally defined mapping $\boldsymbol{x} \mapsto \hat{\boldsymbol{u}}(\boldsymbol{x})$ over the entire domain, capable of being evaluated at arbitrary spatial coordinates. However, being meshless does not imply freedom from geometric considerations: defining the domain, enforcing boundary conditions, and designing effective collocation point distributions remain essential and influential steps.

3. **Hardware Acceleration and Differentiability**: Modern deep learning frameworks are built upon dense linear algebra operations optimized for parallel accelerator hardware (GPUs and TPUs), and provide built-in automatic differentiation for computing spatial derivatives of the represented fields.

3.2.2 Limits of Purely Data-Driven Machine Learning

Despite their high expressivity, deploying purely data-driven black-box networks to fluid mechanics and non-Newtonian rheology reveals fundamental limitations:

- **Data scarcity in experimental rheology**: Deep neural networks typically rely on large labeled datasets. In fluid experiments, acquiring dense velocity fields via Particle Image Velocimetry (PIV) is labor-intensive, while non-intrusive measurements of internal extra-stress components are very hard to obtain in complex geometries.

- **Absence of physical grounding**: Purely data-driven models minimize empirical loss on discrete observations without physical constraints. Consequently, their predictions often violate incompressibility ($\nabla \cdot \boldsymbol{u} \neq 0$), momentum balance, and constitutive laws, generating non-physical artifacts when interpolating or extrapolating.

- **Overfitting**: In sparse-data regimes, over-parameterized networks can easily fit measurement noise, obtaining low training errors while losing accuracy in unmonitored regions.

- **Inability to discover governing parameters**: Black-box models learn input-output mappings, but cannot extract underlying physical properties, such as relaxation times or polymeric viscosities, that govern flow dynamics.

These limitations show that data alone are insufficient for scientific modeling in rheology, motivating frameworks that embed physical conservation laws directly into the loss function.

3.3 Physics-Informed Neural Network

The Physics-Informed Neural Network (PINN) framework, formalized by Raissi, Perdikaris, and Karniadakis [8], addresses the problems of purely data-driven modeling by incorporating differential partial equations directly into the network training objective.

3.3.1 Physics-Constrained Loss Formulation

Following the unified boundary-value problem (BVP) notation established in the scientific machine learning literature [9], consider a physical state described by a vector of continuous fields $\boldsymbol{u}(\boldsymbol{x})$ defined over a spatial domain $\Omega \subset \mathbb{R}^d$ bounded

by $\partial\Omega$:

$$\mathcal{F}[\boldsymbol{u}(\boldsymbol{x});\boldsymbol{\lambda}] = \mathbf{0}, \quad \boldsymbol{x} \in \Omega, \tag{3.9}$$

$$\mathcal{B}[\boldsymbol{u}(\boldsymbol{x})] = \boldsymbol{g}(\boldsymbol{x}), \quad \boldsymbol{x} \in \partial\Omega, \tag{3.10}$$

where:

- $\boldsymbol{u}(\boldsymbol{x})$ denotes the physical state variables (in our viscoelastic framework, $\boldsymbol{u} = (\psi, p, \boldsymbol{\tau})$, comprising stream function, pressure, and extra-stress tensor components);
- $\mathcal{F}$ is a general non-linear differential operator encoding the governing field equations (mass conservation, momentum balance, and constitutive transport equations);
- $\boldsymbol{\lambda}$ denotes physical or constitutive parameters governing the flow dynamics (such as solvent viscosity, polymeric viscosity, and relaxation time);
- $\mathcal{B}$ denotes the boundary operator enforcing kinematic constraints on the boundaries $\partial\Omega$ (such as no-slip conditions along solid walls and prescribed angular velocities on roller surfaces), with prescribed boundary values $\boldsymbol{g}(\boldsymbol{x})$.

In a PINN, the continuous solution fields are parameterized by a deep neural network surrogate $\hat{\boldsymbol{u}}(\boldsymbol{x};\boldsymbol{\theta})$ parameterized by weights and biases $\boldsymbol{\theta}$. Evaluating the governing differential operator on the network outputs via automatic differentiation defines the pointwise differential residual:

$$\boldsymbol{f}(\boldsymbol{x};\boldsymbol{\theta},\boldsymbol{\lambda}) \equiv \mathcal{F}[\hat{\boldsymbol{u}}(\boldsymbol{x};\boldsymbol{\theta});\boldsymbol{\lambda}]. \tag{3.11}$$

The network parameters $\boldsymbol{\theta}$, alongside any unknown physical parameters $\boldsymbol{\lambda}$ in inverse problems, are optimized simultaneously by minimizing a composite multi-

objective loss:

$$\mathcal{L}_{\text{total}}(\boldsymbol{\theta}, \boldsymbol{\lambda}) = w_{\text{data}}\mathcal{L}_{\text{data}} + w_{\text{bc}}\mathcal{L}_{\text{bc}} + w_{\text{pde}}\mathcal{L}_{\text{pde}}, \tag{3.12}$$

where each individual loss component is evaluated as:

$$\mathcal{L}_{\text{data}} = \frac{1}{N_{\text{data}}} \sum_{i=1}^{N_{\text{data}}} \left\| \hat{\boldsymbol{u}}(\boldsymbol{x}_i^{\text{data}}) - \boldsymbol{u}_i^{\text{data}} \right\|_2^2, \tag{3.13}$$

$$\mathcal{L}_{\text{bc}} = \frac{1}{N_{\text{bc}}} \sum_{j=1}^{N_{\text{bc}}} \left\| \mathcal{B}[\hat{\boldsymbol{u}}(\boldsymbol{x}_j^{\text{bc}})] - \boldsymbol{g}(\boldsymbol{x}_j^{\text{bc}}) \right\|_2^2, \tag{3.14}$$

$$\mathcal{L}_{\text{pde}} = \frac{1}{N_{\text{pde}}} \sum_{k=1}^{N_{\text{pde}}} \left\| \boldsymbol{f}(\boldsymbol{x}_k^{\text{pde}}; \boldsymbol{\theta}, \boldsymbol{\lambda}) \right\|_2^2. \tag{3.15}$$

Here, $\{\boldsymbol{x}_i^{\text{data}}\}$ denotes available empirical observations or measurement locations, $\{\boldsymbol{x}_j^{\text{bc}}\}$ represents points sampled along physical boundaries $\partial\Omega$, and $\{\boldsymbol{x}_k^{\text{pde}}\}$ designates interior *collocation points* scattered across Ω.

This formulation naturally supports two complementary problem settings:

- **Direct problems**: The physical parameters $\boldsymbol{\lambda}$ are known. The network optimizes its weights $\boldsymbol{\theta}$ to solve for the continuous fields $\hat{\boldsymbol{u}}(\boldsymbol{x})$, acting as a meshless continuous PDE solver.

- **Inverse problems**: One or more constitutive parameters in $\boldsymbol{\lambda}$ are unknown. The network simultaneously reconstructs the unobserved latent fields and identifies the unknown parameters by combining empirical velocity observations $\mathcal{L}_{\text{data}}$ with physical consistency constraints $\mathcal{L}_{\text{pde}}$.

A key methodological principle of this thesis is the distinction between field reconstruction accuracy and parameter identifiability. A network may fit observed velocity fields accurately, yet inferred rheological parameters may remain uncertain if the kinematics lack sufficient variation or if the loss landscape exhibits degeneracies. Embedding physical equations provides essential structure,

but does not automatically guarantee unique parameter recovery.

3.3.2 PINN-Specific Features

Several operational features distinguish physics-informed training from conventional supervised learning:

1. **Training as a boundary-value solver**: In classical machine learning, a model is trained on a sample set to generalize to unseen statistical distributions. In PINNs, *training is the numerical solver*. The network is not designed to predict flows across arbitrary unseen geometries; rather, optimization directly solves the specific problem on the specified domain Ω. Available measurements $\mathcal{L}_{\text{data}}$ act as an active data anchor that guides the network toward the true physical state.

2. **Physical regularization and dynamic loss weighting**: In standard supervised training, the network is unconstrained between data points. In PINNs, evaluating differential residuals $\mathcal{L}_{\text{pde}}$ across interior collocation points penalizes non-physical gradients and curvatures, restricting the admissible solution space. However, balancing the competing loss weights (w_{data}, w_{bc}, and w_{pde}) presents practical challenges. As analyzed by Wang, Teng, and Perdikaris [23], the backpropagated gradient magnitudes arising from PDE residuals and boundary constraints often differ by several orders of magnitude. This imbalance induces numerical stiffness, where boundary terms can dominate and suppress the learning signal of interior PDE residuals (or vice versa). Calibrated weighting between data, boundary, and differential components is therefore needed to prevent optimization stalls and ensure stable convergence.

3. **Coordinate Automatic Differentiation (AD)**: Automatic differentia-

tion serves a dual purpose in PINNs: beyond computing loss gradients with respect to network parameters, it evaluates exact spatial derivatives of the predicted physical fields with respect to coordinates:

$$\text{Coordinate AD (PINNs):} \quad \frac{\partial \hat{\boldsymbol{u}}}{\partial \boldsymbol{x}}, \quad \frac{\partial^2 \hat{\boldsymbol{u}}}{\partial \boldsymbol{x}^2}. \tag{3.16}$$

The accuracy of these derivatives with respect to the true physical field depends on how closely the network itself approximates the physical solution.

4. **Collocation point sampling**: Unlike traditional CFD solvers that require meshes and element connectivity tables, PINNs evaluate differential residuals at sparse collocation points $\{\boldsymbol{x}_k^{\text{pde}}\}$. This meshless feature simplifies domain representation, though sufficient point density remains necessary in boundary layers and near stagnation regions.

Table 3.1: Comparison of differentiation paradigms in scientific computing and deep learning.

Method	Formulation	Truncation Error	Grid Dependence
Finite Differences (FDM)	Discrete difference quotients (Taylor series)	$\mathcal{O}(\Delta x^2)$	Structured Mesh Stencil
Symbolic Differentiation	Closed-form analytical expression	None	Expression Growth
Automatic Differentiation (AD)	Exact chain rule on computational graph	**None (Machine Precision)**	**Meshless / Continuous**

3.4 Foundational Literature Benchmarks

To contextualize the development of physics-informed methods in fluid mechanics and non-Newtonian rheology, we review two foundational papers that directly guide the architectural and physical formulations adopted in this work: the incompressible Navier–Stokes benchmark established by Raissi, Perdikaris, and Karniadakis [8], and the *ViscoelasticNet* framework formulated by Thakur, Raissi, and Ardekani [11].

3.4.1 Navier–Stokes Flow Past a Cylinder

In their foundational contribution, Raissi et al. [8] demonstrated that Physics-Informed Neural Networks could simultaneously solve unmeasured latent fields and discover governing hydrodynamic transport coefficients in unsteady, two-dimensional flow past a cylinder at Reynolds number $Re = 100$.

The dimensionless governing momentum and continuity equations were formulated as:

$$u_t + \lambda_1(uu_x + vu_y) = -p_x + \lambda_2(u_{xx} + u_{yy}), \tag{3.17}$$

$$v_t + \lambda_1(uv_x + vv_y) = -p_y + \lambda_2(v_{xx} + v_{yy}), \tag{3.18}$$

$$u_x + v_y = 0, \tag{3.19}$$

where $\lambda_1 = 1.0$ governs convective acceleration and $\lambda_2 = 1/Re = 0.01$ represents the dimensionless kinematic viscosity.

Instead of penalizing continuity divergence as an empirical loss objective, the authors parameterised the two-dimensional velocity field through a scalar stream function $\psi(t, x, y)$ via $u = \partial\psi/\partial y$ and $v = -\partial\psi/\partial x$. By construction, the stream function satisfies the continuity constraint (3.19) identically across the domain, reducing competition between mass and momentum balance. A deep feedforward architecture comprising 9 hidden layers with 20 neurons per layer, activated by hyperbolic tangent functions, was trained to output (ψ, p) while treating λ_1 and λ_2 as dynamically learnable scalar parameters.

When trained on spatio-temporal velocity observations sampled behind the cylinder, the framework revealed two main capabilities. First, even in the complete absence of pressure measurements in the flow domain, the network reconstructed the unobserved continuous pressure field $p(t, x, y)$ exclusively from kinematic velocity data and momentum balance, up to an arbitrary constant. Second,

the solver demonstrated robustness to Gaussian noise: on clean data, the parameters converged within 0.08% error for λ_1 and 4.67% for λ_2, and remained strictly bounded within 0.17% and 5.70% error even when corrupted by 1% uncorrelated Gaussian noise.

From a continuum rheology perspective, however, a critical limitation characterizes the Navier-Stokes cylinder benchmark, considering the complete absence of any information about stresses and elasticity.

3.4.2 ViscoelasticNet Framework

Recognizing this fundamental distinction, Thakur, Raissi, and Ardekani [11] first extended physics-informed neural networks to non-Newtonian fluid mechanics through the *ViscoelasticNet* framework. In polymeric liquids, fluid elements retain fading memory of past conformations, requiring the momentum balance to be coupled with tensorial, frame-indifferent constitutive transport equations governed by the Upper-Convected Time Derivative (UCTD).

ViscoelasticNet addressed two interconnected bottlenecks that severely hamper traditional computational rheology. On the experimental side, non-intrusive full-field measurements of internal polymeric extra-stress components $\boldsymbol{\tau}$ are generally inaccessible in complex flow geometries. On the numerical simulation side, mesh-based discretization methods (FEM and FVM) suffer from the High-Weissenberg Number Problem (HWNP).

To resolve these difficulties, Thakur et al. proposed an inverse physics-informed paradigm operating as a virtual rheometer. By embedding a unified constitutive expression of Oldroyd-B, Giesekus, and linear Phan-Thien–Tanner (PTT) model:

$$\boldsymbol{\tau} + \lambda \overset{\nabla}{\boldsymbol{\tau}} + \frac{\alpha\lambda}{\mu_p}(\boldsymbol{\tau}\cdot\boldsymbol{\tau}) + \frac{\varepsilon\lambda}{\mu_p}\mathrm{tr}(\boldsymbol{\tau})\boldsymbol{\tau} = 2\mu_p \boldsymbol{D}, \tag{3.20}$$

the network was capable to work with different fluids nature, while simultaneously

infer unmeasured extra-stress components $(\tau_{xx}, \tau_{xy}, \tau_{yy})$ and discover unknown material parameters $(\lambda, \mu_p, \alpha, \varepsilon)$.

To overcome the numerical stiffness arising from the mathematical character of the coupled equations, ViscoelasticNet introduced a decoupled two-stage training sequence. In Stage 1, the velocity and stress networks were optimized together by minimizing the sum of kinematic data loss, boundary penalties, and constitutive residuals, in order to discover the rheological parameters $(\lambda, \mu_p, \alpha, \varepsilon)$ and find the stress field. In Stage 2, the learned stress network was frozen, while the pressure network and solvent viscosity μ_s were trained against the momentum balance equations and the velocity observation.

ViscoelasticNet provided proof-of-concept evidence that neural networks can reconstruct non-Newtonian stress fields, reporting relative L_2 stress errors below 5% and robust convergence under synthetic velocity noise up to 10%. Anyway, a critical examination of its methodological scope reveals an essential research gap. ViscoelasticNet was evaluated exclusively on open, pressure-driven geometries, a planar stenosed channel and a symmetric cross-slot configuration. Both benchmark domains are characterized by well-defined inflow and outflow boundaries, continuous net mass throughput, and a dominant convective pressure drop that provides strong directional alignment between velocity and pressure gradients.

More broadly, extending this physics-informed framework to increasingly complex geometries represents an essential step to validate its robustness, testing whether neural rheometers can reliably handle realistic and challenging flow configurations.

Chapter 4

Numerical Results and Methodology

The previous chapters established the governing equations of viscoelastic flow (Chapter 2) and the physics-informed neural network framework (Chapter 3). In this chapter, we evaluate that framework on the four-roll mill benchmark. The goal is double: first, to test direct field reconstruction against high-fidelity numerical solutions; second, to assess the inverse PINN as a virtual rheometer capable of reconstructing hidden stress and pressure fields while extracting unknown constitutive parameters directly from kinematic observations.

We first present the problem formulation in the recirculating cavity, the reference finite element dataset generated with COMSOL Multiphysics, and the decoupled network architecture. We then evaluate direct field reconstruction, report inverse discovery results across Oldroyd-B, Giesekus, and PTT fluids, and examine the physical limits, invariances, and degeneracies that govern viscoelastic parameter estimation.

4.1 Problem Formulation

4.1.1 Recirculating Cavity Challenges

Unlike conventional benchmarks in open channels with net through-flow (such as planar contractions or cross-slots), the four-roll mill [16] is a closed recirculating cavity driven only by four rotating cylinders with zero net mass flux.

This geometry introduces some computational and physical challenges:

- **Mixed Flow Regimes**: Within the same compact domain, the flow could couple intense simple shear in the narrow gaps near the cylinders with an extensional stagnation regime at the cavity center. In this work, we focus specifically on this configuration, that generates a pure planar extensional flow at the central stagnation point.

- **Stress Recirculation along Closed Streamlines**: Polymeric extra-stresses along closed streamlines and without open inlets or outlets can be hard to analyze, requiring the network to capture steep localized stress layers.

4.1.2 Direct and Inverse Modes

The computational framework operates in two distinct modes:

- **Direct Mode**: Direct Mode, also called *forward problem* in the scientific machine learning and PINN literature, denotes the reconstruction of unobserved stress and pressure fields $(\boldsymbol{\tau}, p)$ when the constitutive parameters $(\lambda, \mu_p, \mu_s, \alpha, \varepsilon)$ are known. Planar velocity observations (representative of experimental PIV data) supervise the stream function ψ in the interior alongside boundary conditions, while the network reconstructs the unmeasured extra-stress tensor $\boldsymbol{\tau}(x, y)$ and hydrodynamic pressure $p(x, y)$ using

the momentum and constitutive residuals without interior stress or pressure data.

- **Inverse Mode (Virtual Rheometer)**: Kinematic velocity observations within the cavity are supplied together with roller boundary values. The network reconstructs the unobserved latent stress and pressure fields while simultaneously discovering the governing physical parameters $(\lambda, \mu_p, \mu_s, \alpha, \varepsilon)$ through physics-constrained optimization.

Table 4.1 summarizes the role of each field variable and parameter across both modes.

Table 4.1: Classification of physical fields and parameters across direct and inverse problem formulations.

Physical Quantity	**Symbol**	**Direct Problem**	**Inverse Problem**
Velocity Field	$\boldsymbol{u} = (u, v)$	Observed (via ψ)	Observed (via ψ)
Pressure Field	p	Latent unknown (PDE-resolved)	Latent unknown (PDE-resolved)
Polymeric Extra-Stress	$\boldsymbol{\tau} = (\tau_{xx}, \tau_{xy}, \tau_{yy})$	Latent unknown (PDE-resolved)	Latent unknown (PDE-resolved)
Relaxation Time	λ (Wi)	Prescribed known	Unknown parameter
Polymeric Viscosity	μ_p	Prescribed known	Unknown parameter
Solvent Viscosity	μ_s	Prescribed known	Unknown parameter
Giesekus Mobility	α	Prescribed known	Unknown parameter
PTT Extensibility	ε	Prescribed known	Unknown parameter

4.2 Reference CFD Dataset

4.2.1 Finite Element Setup

Reference numerical solutions were generated using COMSOL Multiphysics [12], solving steady incompressible viscoelastic flow on unstructured triangular meshes. The governing momentum and continuity equations were discretized using finite elements method. To ensure that reference fields are free from discretization artifacts, we evaluated spatial mesh convergence.

The canonical Oldroyd-B benchmark at $\lambda = 1.0\,\mathrm{s}$ $(Wi = 1.667)$ was selected as the representative test case. Because linear Oldroyd-B lacks quadratic or

exponential stress saturation mechanisms, normal stresses grow steeply near the central stagnation point, providing a conservative worst-case baseline for spatial resolution. Mesh independence was verified on this representative configuration and confirmed across all other fluid models and regimes investigated.

We evaluated six progressive grid resolutions:

- **M5k (Coarse Grid)**: 5,086 nodes, capturing macroscopic cavity recirculation.
- **M12k (Light Grid)**: 12,760 nodes.
- **M29k (Intermediate Grid)**: 29,392 nodes.
- **M52k (Fine Grid)**: 52,648 nodes.
- **M88k (Very Fine Grid)**: 88,804 nodes.
- **M125k (Dense Baseline)**: 125,456 nodes, resolving sub-millimeter viscoelastic stress layers within the narrow roller gaps.

Figure 4.1 illustrates the triangular mesh and spatial distribution of the $\mathcal{N}_{\text{coll}} =$ 5,086 collocation points for the coarse baseline grid (M5k), highlighting the higher point density concentrated near the cylinder surfaces and within the narrow roller gaps.

4.2.2 Mesh Convergence Study

To quantitatively verify local kinematic and stress fields across mesh densities, three diagnostic cutlines are established (Figure 4.2):

1. **Horizontal Cutline** ($y = 0$, $-L/2 \leq x \leq L/2$): Outflow elongational axis passing through the central hyperbolic stagnation point $(0, 0)$, along which tensile normal stress tails τ_{xx} develop.

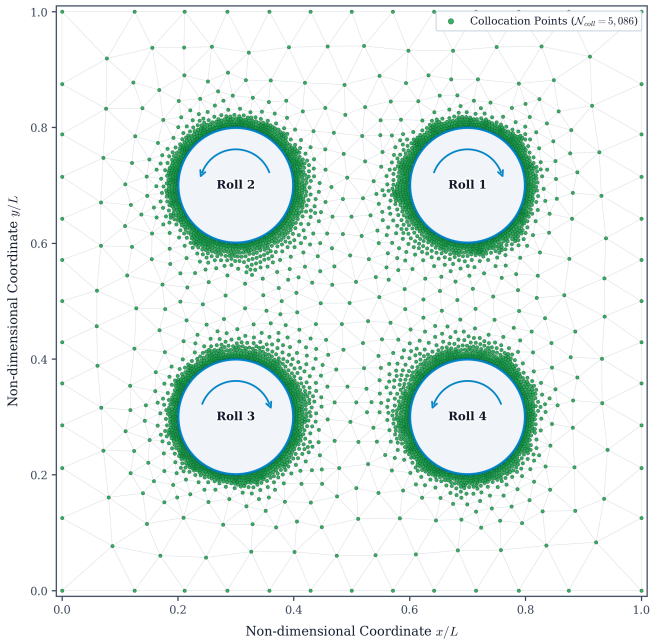


Figure 4.1: Spatial distribution of the $\mathcal{N}_{\text{coll}} = 5{,}086$ interior collocation points and underlying unstructured triangular mesh for the coarse resolution grid (M5k), showing the local point refinement around the four rotating cylinder boundaries and within the roller gaps.

2. **Vertical Cutline** ($x = 0, -L/2 \leq y \leq L/2$): Compressive inflow axis capturing fluid deceleration toward the stagnation zone.

3. **Oblique Diagonal Cutline** ($y = x$): Connects opposite cavity corners across roller nips, capturing the transition from the extensional core into high-shear roller boundary layers.

Figure 4.3 displays spatial profiles of velocity and extra-stress tensor components $(u, v, \tau_{xx}, \tau_{xy}, \tau_{yy})$ across the cutlines. Spatial velocity profiles between M52k, M88k, and M125k deviate by less than 0.05%, while peak shear stresses in roller gaps differ by under 0.3%, confirming that M125k is fully grid-independent and provides a reliable ground truth.

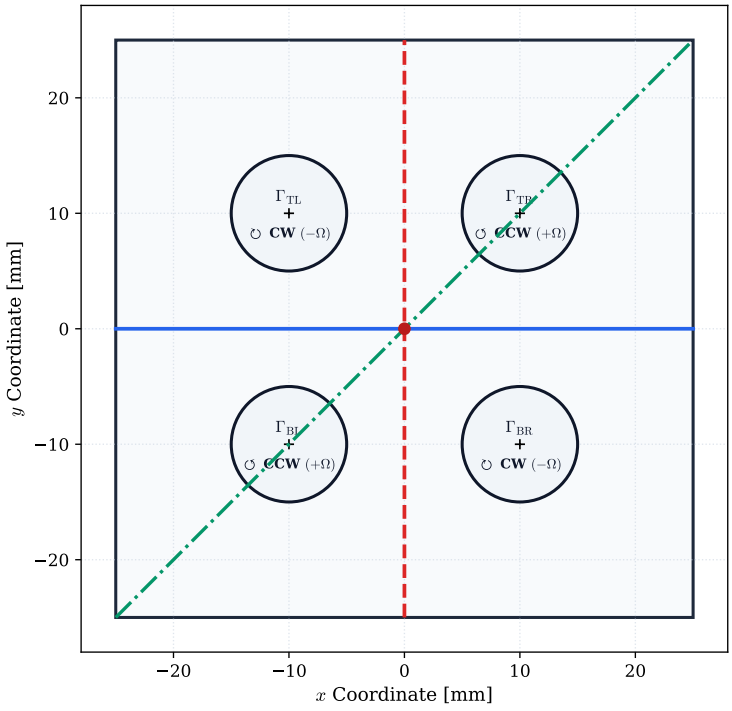


Figure 4.2: Schematic of the four-roll mill geometry with the four counter-rotating cylinders, central extensional stagnation point $(0, 0)$ (red circle), and three diagnostic verification cutlines: Horizontal ($y = 0$, outflow, blue solid), Vertical ($x = 0$, inflow, red dashed), and Oblique Diagonal ($y = x$, roller nips, green dash-dot).

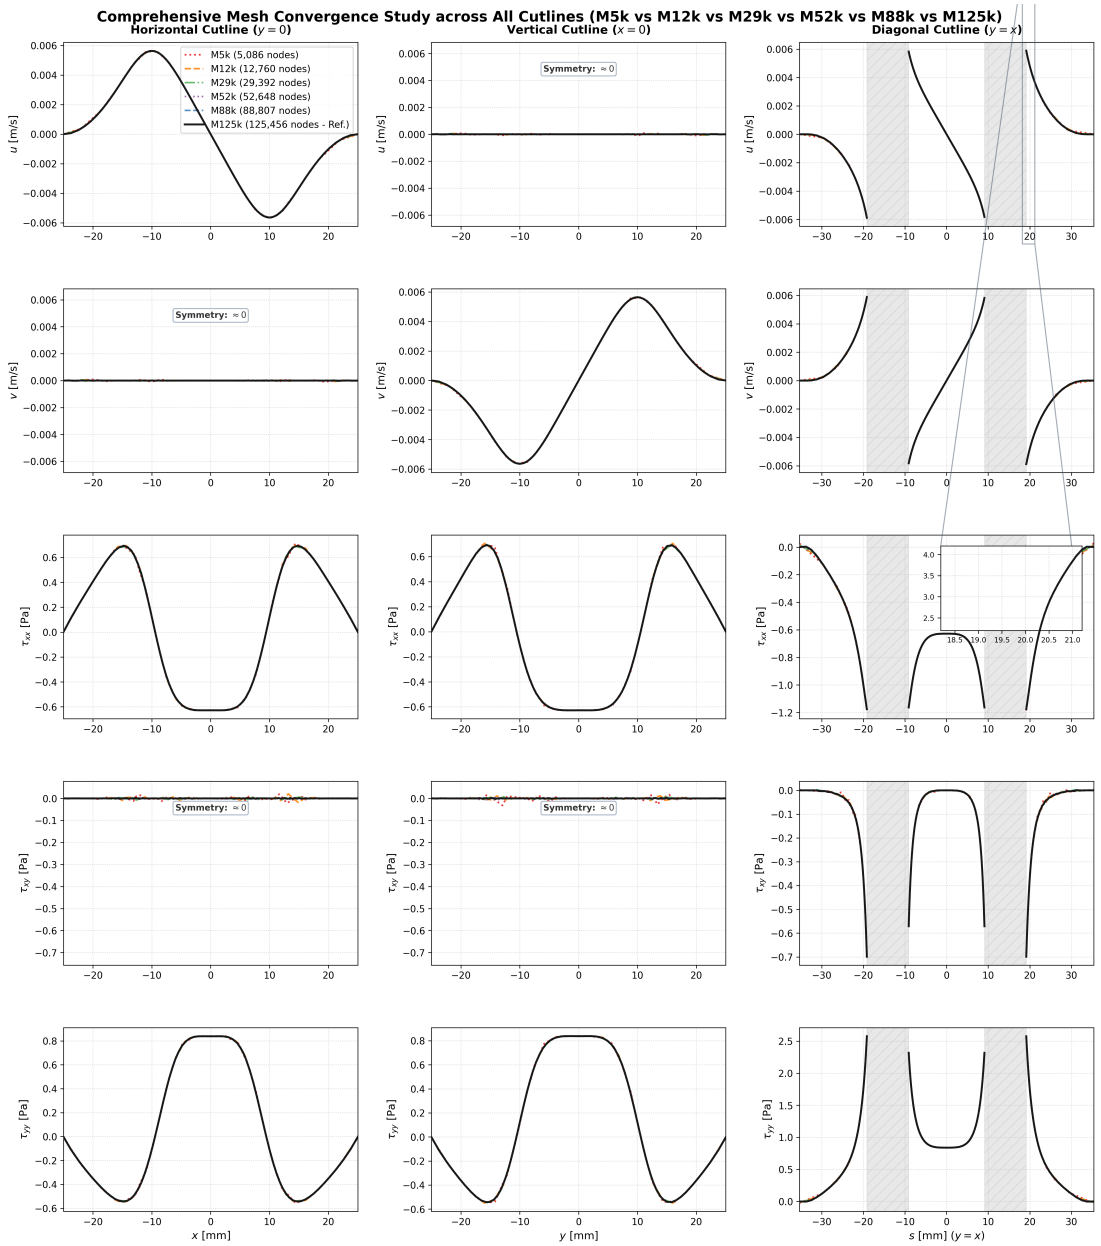


Figure 4.3: High-fidelity COMSOL mesh convergence matrix for the representative Oldroyd-B fluid ($\lambda = 1.0\,\mathrm{s}$, $\mu_p = 0.50\,\mathrm{Pa}\cdot\mathrm{s}$, $\mu_s = 0.50\,\mathrm{Pa}\cdot\mathrm{s}$, $Wi = 1.667$, $Re = 0.0417$, $\alpha = 0$, $\varepsilon = 0$) across six mesh densities (M5k to M125k) along the Horizontal ($y = 0$), Vertical ($x = 0$), and Oblique Diagonal ($y = x$) cutlines for velocity components (u, v) and stress tensor components $(\tau_{xx}, \tau_{xy}, \tau_{yy})$. Shaded gray vertical bands in the diagonal cutline panels indicate the solid interior occupied by the cylindrical rollers.

4.3 Network Architecture and Loss Formulation

4.3.1 Decoupled Three-Head Network

Following the stream function formulation introduced in Section 2.2.1 (Eq. (2.10)), the velocity field identically satisfies mass conservation ($\nabla \cdot \boldsymbol{u} \equiv 0$). As shown in Figure 4.4, the physical fields are predicted by three independent Multi-Layer Perceptron (MLP) branches rather than a single shared network:

$$(x, y) \longrightarrow \begin{cases} \psi(x, y) & \text{(Stream function, yielding } u, v \text{ via AD)}, \\ p(x, y) & \text{(Hydrodynamic pressure)}, \\ \boldsymbol{\tau}(x, y) = (\tau_{xx}, \tau_{xy}, \tau_{yy}) & \text{(Extra-stress tensor)}. \end{cases} \tag{4.1}$$

Each branch contains 8 hidden layers with 128 neurons per layer, activated by the Sigmoid Linear Unit (SiLU$(z) = z/(1+e^{-z})$). Its smooth, infinitely differentiable profile provides well-behaved second and third-order derivative graphs during automatic differentiation.

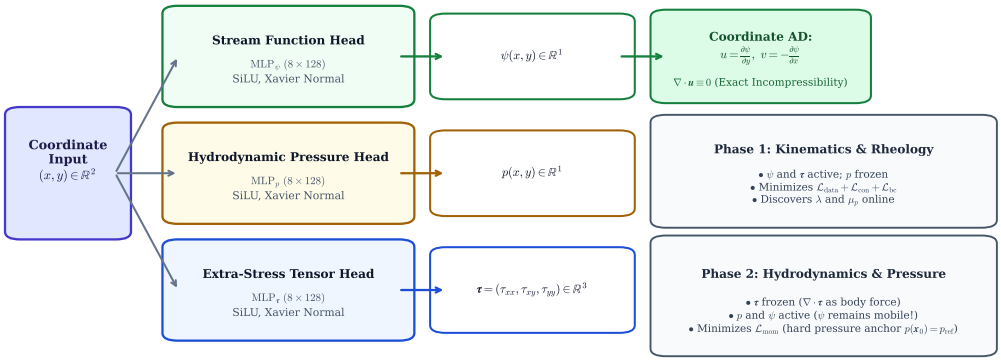


Figure 4.4: Decoupled three-head neural network architecture for the viscoelastic four-roll mill solver, showing coordinate inputs, separate MLP heads for stream function ψ, pressure p, and stress tensor $\boldsymbol{\tau}$, coordinate automatic differentiation, and the two-stage optimization schedule.

4.3.2 Weight Initialization

Hidden layers are initialized with the Xavier normal distribution, scaling weight variance inversely with layer width to maintain balanced gradient amplitudes across depth.

The final linear layers of the extra-stress network $\mathcal{N}_{\boldsymbol{\tau}}$ and pressure network $\mathcal{N}_p$ are initialized to zero ($\boldsymbol{W}^{(L)} = \mathbf{0}, \boldsymbol{b}^{(L)} = \mathbf{0}$). This choice provides two practical benefits:

1. **Better Starting State**: At iteration zero, the model predicts zero extra-stress ($\boldsymbol{\tau} = \mathbf{0}$) and uniform pressure, reflecting a fluid initially at rest before deformation begins.

2. **Numerical Stability**: Random initial weights in the stress network produce high-frequency spatial fluctuations. Because constitutive models contain nonlinear terms ($\alpha\boldsymbol{\tau}^2$ in Giesekus) and convective terms ($\boldsymbol{u} \cdot \nabla\boldsymbol{\tau}$), uncalibrated stresses at step zero can generate extreme loss spikes ($> 10^6$) that disrupt the optimizer. Initializing output layers to zero avoids these early instabilities.

4.3.3 Governing Residuals and Scaling

Following the nondimensionalization introduced in Section 2.4, characteristic scales are $H_{\text{ref}} = R = 0.005\,\text{m}$ (roller radius) and $U_{\text{ref}} = u_{\text{roll}} = 8.333 \times 10^{-3}\,\text{m/s}$. In the inverse problem, fluid viscosities μ_s and μ_p are unknown a priori. Scaling differential residuals by evolving parameter estimates would introduce severe numerical feedback and destabilize backpropagation. To decouple computational normalization from evolving parameter values, the solver employs a fixed reference dynamic viscosity scale $\eta_0 = 1.0\,\text{Pa} \cdot \text{s}$, alongside the roller radius

$H_{\text{ref}} = R = 0.005\,\text{m}$ and peripheral roller velocity $U_{\text{ref}} = u_{\text{roll}} = 8.333 \times 10^{-3}\,\text{m/s}$:

$$\mu_s^* = \frac{\mu_s}{\eta_0}, \quad \mu_p^* = \frac{\mu_p}{\eta_0}, \quad Re_{\text{scale}} = \frac{\rho U_{\text{ref}} H_{\text{ref}}}{\eta_0}, \quad Wi = \frac{\lambda U_{\text{ref}}}{H_{\text{ref}}}. \tag{4.2}$$

Spatial coordinates are mapped to the unit box $[0,1]^2$ using the cavity width $H_{\text{coord}} = L = 0.05\,\text{m}$ to keep activation arguments well conditioned. Derivative operators with respect to physical space are scaled accordingly via Automatic Differentiation by $(H_{\text{ref}}/H_{\text{coord}})$.

In contrast to the coordinates, the physical output fields (p and stress components $\tau_{xx}, \tau_{xy}, \tau_{yy}$) oscillate symmetrically with both positive and negative values across the cavity. To keep raw network outputs in the balanced range $[-1,1]$, each output head is scaled by its maximum absolute reference value $(p_{\text{scale}}, \boldsymbol{\tau}_{\text{scale}})$:

$$p(x,y) = p_{\text{scale}} \cdot \hat{p}(x,y), \qquad \boldsymbol{\tau}(x,y) = \boldsymbol{\tau}_{\text{scale}} \cdot \hat{\boldsymbol{\tau}}(x,y), \tag{4.3}$$

where $\hat{p}, \hat{\boldsymbol{\tau}} \in [-1,1]$ denote the normalized output representations. Denoting the scaled physical fields evaluated at interior collocation points $\boldsymbol{x}_m \in \Omega$ simply as $(u, v, p, \boldsymbol{\tau})$, the governing PDE residuals evaluate explicitly to:

1. **Linear Momentum Residuals** (f_u, f_v):

$$f_u = Re_{\text{scale}} \left(u\frac{\partial u}{\partial x} + v\frac{\partial u}{\partial y} \right) + \frac{\partial p}{\partial x} - \mu_s^* \left(\frac{\partial^2 u}{\partial x^2} + \frac{\partial^2 u}{\partial y^2} \right) - \left(\frac{\partial \tau_{xx}}{\partial x} + \frac{\partial \tau_{xy}}{\partial y} \right), \tag{4.4}$$

$$f_v = Re_{\text{scale}} \left(u\frac{\partial v}{\partial x} + v\frac{\partial v}{\partial y} \right) + \frac{\partial p}{\partial y} - \mu_s^* \left(\frac{\partial^2 v}{\partial x^2} + \frac{\partial^2 v}{\partial y^2} \right) - \left(\frac{\partial \tau_{xy}}{\partial x} + \frac{\partial \tau_{yy}}{\partial y} \right). \tag{4.5}$$

2. **Viscoelastic Constitutive Residuals** ($f_{\tau_{xx}}, f_{\tau_{xy}}, f_{\tau_{yy}}$):

$$f_{\tau_{xx}} = \mathcal{G}_{xx}(\boldsymbol{\tau}) + Wi\,\overset{\nabla}{\tau}_{xx} - 2\mu_p^*\frac{\partial u}{\partial x}, \tag{4.6}$$

$$f_{\tau_{yy}} = \mathcal{G}_{yy}(\boldsymbol{\tau}) + Wi\,\overset{\nabla}{\tau}_{yy} - 2\mu_p^*\frac{\partial v}{\partial y}, \tag{4.7}$$

$$f_{\tau_{xy}} = \mathcal{G}_{xy}(\boldsymbol{\tau}) + Wi\,\overset{\nabla}{\tau}_{xy} - \mu_p^*\left(\frac{\partial u}{\partial y} + \frac{\partial v}{\partial x}\right), \tag{4.8}$$

where $\overset{\nabla}{\tau}_{ij}$ denotes the upper-convected derivative (Eq. (2.17)), and $\boldsymbol{\mathcal{G}}(\boldsymbol{\tau})$ defines the constitutive model:

$$\boldsymbol{\mathcal{G}}(\boldsymbol{\tau}) = \begin{cases} \boldsymbol{\tau} & \text{(Oldroyd-B)}, \\ \boldsymbol{\tau} + \dfrac{\alpha Wi}{\mu_p^*}\boldsymbol{\tau}^2 & \text{(Giesekus)}, \\ \left(1 + \dfrac{\varepsilon Wi}{\mu_p^*}\mathrm{tr}(\boldsymbol{\tau})\right)\boldsymbol{\tau} & \text{(Linear PTT)}. \end{cases} \tag{4.9}$$

4.3.4 Loss Formulation

The composite objective function minimized during training is:

$$\mathcal{L}_{\text{total}} = w_{\text{data}}\mathcal{L}_{\text{data}} + w_{\text{bc}}\mathcal{L}_{\text{bc}} + w_{\text{con}}\mathcal{L}_{\text{con}} + w_{\text{mom}}\mathcal{L}_{\text{mom}} + w_{\text{roll_stress}}\mathcal{L}_{\text{roll_stress}}. \tag{4.10}$$

Because spatial coordinates ($[0,1]^2$), kinematic velocities ($\mathcal{O}(1)$ via U_{ref} and variance scaling), and stress/pressure outputs are systematically scaled to dimensionless order-unity values, the individual loss components remain naturally balanced without requiring stiff empirical weighting. Accordingly, the baseline weights are set uniformly to unity: $w_{\text{data}} = 1.0$, $w_{\text{bc}} = 1.0$, and $w_{\text{con}} = 1.0$. The momentum balance weight follows the staged training schedule ($w_{\text{mom}} = 0.0$ in Phase 1, $w_{\text{mom}} = 1.0$ in Phase 2), while $w_{\text{roll_stress}} = 1.0$ is selectively activated during inverse parameter discovery and disabled ($w_{\text{roll_stress}} = 0$) during direct field reconstruction.

The individual loss terms are defined as follows:

- **Kinematic Data Loss** ($\mathcal{L}_{\text{data}}$): Quantifies the mean squared error against observed dimensionless velocity data:

$$\mathcal{L}_{\text{data}} = \frac{1}{N_{\text{data}}} \sum_{i=1}^{N_{\text{data}}} \left[\frac{(\hat{u}(\boldsymbol{x}_i) - u_i^{\text{data}})^2}{\sigma_u^2} + \frac{(\hat{v}(\boldsymbol{x}_i) - v_i^{\text{data}})^2}{\sigma_v^2} \right], \tag{4.11}$$

  where $\hat{u}(\boldsymbol{x}) = \frac{\partial \hat{\psi}}{\partial y}$, $\hat{v}(\boldsymbol{x}) = -\frac{\partial \hat{\psi}}{\partial x}$, and σ_u^2, σ_v^2 are empirical dimensionless field variances used to scale horizontal and vertical kinematic errors to comparable magnitude.

- **Kinematic Boundary Loss** ($\mathcal{L}_{\text{bc}}$): Enforces no-slip conditions on outer cavity walls and prescribed tangential velocities on the rotating cylinder surfaces in dimensionless variables:

$$\mathcal{L}_{\text{bc}} = \frac{1}{N_{\text{bc}}} \sum_{j=1}^{N_{\text{bc}}} \left\| \hat{\boldsymbol{u}}(\boldsymbol{x}_j) - \boldsymbol{u}^{\text{wall/roll}}(\boldsymbol{x}_j) \right\|_2^2. \tag{4.12}$$

- **Constitutive Residual Loss** ($\mathcal{L}_{\text{con}}$): Evaluates the mean squared residual of the rheological equations over collocation points:

$$\mathcal{L}_{\text{con}} = \frac{1}{N_{\text{col}}} \sum_{m=1}^{N_{\text{col}}} \left[f_{\tau_{xx}}(\boldsymbol{x}_m)^2 + f_{\tau_{xy}}(\boldsymbol{x}_m)^2 + f_{\tau_{yy}}(\boldsymbol{x}_m)^2 \right]. \tag{4.13}$$

- **Momentum Residual Loss** ($\mathcal{L}_{\text{mom}}$): Evaluates linear momentum balance over collocation points:

$$\mathcal{L}_{\text{mom}} = \frac{1}{N_{\text{col}}} \sum_{m=1}^{N_{\text{col}}} \left[f_u(\boldsymbol{x}_m)^2 + f_v(\boldsymbol{x}_m)^2 \right]. \tag{4.14}$$

- **Pressure Fixing**: In incompressible flow with enclosed boundaries, the momentum equations determine pressure only through its spatial gradient

∇p. Consequently, the absolute hydrostatic level is mathematically indeterminate up to an arbitrary additive constant. Rather than imposing a soft penalty loss with an empirical weight w_p, the pressure sub-network $\mathcal{N}_p$ is constrained by an exact algebraic hard anchor directly within the forward pass:

$$p(\boldsymbol{x}) = p_{\text{scale}} \left(\mathcal{N}_p(\boldsymbol{x}) - \mathcal{N}_p(\boldsymbol{x}_0) \right) + p_{\text{ref}}, \tag{4.15}$$

where $\boldsymbol{x}_0$ denotes a prescribed fixed boundary reference point (selected on the external wall) and p_{ref} is the corresponding reference Dirichlet value. This construction identically guarantees $p(\boldsymbol{x}_0) \equiv p_{\text{ref}}$ for all network parameters $\boldsymbol{\theta}_p$, acting as an exact single-point Dirichlet gauge without introducing competing loss terms or gradient conflicts.

4.3.5 Roller Stress Supervision

In inverse discovery mode, extra-stress observations are restricted exclusively to the rotating cylinder boundaries Γ_{roll}:

$$\mathcal{L}_{\text{roll_stress}} = \frac{1}{N_{\text{rolls}}} \sum_{k=1}^{N_{\text{rolls}}} \left\| \boldsymbol{\tau}_k - \boldsymbol{\tau}_k^{\text{ref}} \right\|_2^2 . \tag{4.16}$$

This localized boundary condition provides a necessary reference to anchor the stress scale and couple kinematic fields with the constitutive relations during parameter inversion, without requiring any stress measurements inside the fluid domain. In the direct problem, where fluid parameters are prescribed, this supervision term is disabled ($w_{\text{roll_stress}} = 0$).

4.4 Training Strategy

4.4.1 Two-Stage Physical Decoupling

Optimizing momentum and constitutive equations together from scratch often leads to numerical instability, causing the optimizer to stall or produce unphysical fields.

To avoid this imbalance, we split training into two sequential stages:

1. **Phase 1: Kinematics and Rheology**: The momentum balance is turned off ($w_{\text{mom}} = 0$) and the pressure network $\mathcal{N}_p$ remains frozen. The stream function ψ and extra-stress $\boldsymbol{\tau}$ are trained jointly by minimizing kinematic data matching alongside constitutive residuals:

$$\mathcal{L}_{\text{Phase 1}} = w_{\text{data}}\mathcal{L}_{\text{data}} + w_{\text{bc}}\mathcal{L}_{\text{bc}} + w_{\text{con}}\mathcal{L}_{\text{con}} + w_{\text{roll_stress}}\mathcal{L}_{\text{roll_stress}}. \tag{4.17}$$

   In inverse mode, unknown fluid parameters (λ, μ_p, α) are identified during this phase, isolating constitutive dynamics from hydrodynamic pressure interactions.

2. **Phase 2: Hydrodynamics and Pressure Recovery**: The extra-stress network $\mathcal{N}_{\boldsymbol{\tau}}$ is frozen, and its spatial divergence $\nabla \cdot \boldsymbol{\tau}$ is fixed. The momentum loss is activated ($w_{\text{mom}} = 1.0$):

$$\mathcal{L}_{\text{Phase 2}} = w_{\text{data}}\mathcal{L}_{\text{data}} + w_{\text{bc}}\mathcal{L}_{\text{bc}} + w_{\text{mom}}\mathcal{L}_{\text{mom}}. \tag{4.18}$$

Crucially, the stream function network $\mathcal{N}_\psi$ remains trainable alongside $\mathcal{N}_p$ during Phase 2. Keeping ψ mobile allows the kinematic field to make minor mutual adjustments that reconcile viscous and elastic forces with the pressure gradient ∇p, avoiding the numerical stagnation that occurs if the velocity field is

rigidly frozen after Phase 1.

Table 4.2 summarizes the training stages, active heads, and loss weights.

Table 4.2: Summary of the decoupled two-stage training strategy.

Stage	Trainable Heads	Frozen Heads	Active Objectives
Phase 1	ψ, $\boldsymbol{\tau}$	p	$\mathcal{L}_{\text{data}} + \mathcal{L}_{\text{bc}} + \mathcal{L}_{\text{con}} + \mathcal{L}_{\text{roll_stress}}$
Phase 2	p, ψ (mobile)	$\boldsymbol{\tau}$	$\mathcal{L}_{\text{data}} + \mathcal{L}_{\text{bc}} + \mathcal{L}_{\text{mom}}$

4.4.2 Dual-Precision Optimization

Optimization proceeds in two sequential precision tiers:

1. **Tier 1: Fast Exploration via Adam (FP32)**: In early iterations, the loss landscape contains local minima and steep gradients. The Adam optimizer runs in single precision (FP32) to take advantage of GPU throughput. A cosine annealing schedule ($2.5 \times 10^{-3} \rightarrow 2.5 \times 10^{-6}$) guides parameter values into the attraction basin of the physical solution.

2. **Tier 2: Fine-Tuning with L-BFGS in Double Precision (FP64)**: To resolve higher-order spatial derivatives without round-off errors and eliminate stochastic fluctuations, the model switches to double precision (`torch.float64`). L-BFGS with a Strong Wolfe line search and history size $m = 50$ provides superlinear convergence, driving PDE residuals down to tight numerical tolerances.

4.4.3 Transfer Learning

Training PINNs from scratch for every fluid regime requires substantial computational budgets (50,000 iterations) and risks local minima at high elasticity. We establish a Transfer Learning protocol:

1. Pre-train the base network on a moderate baseline case ($\lambda = 0.05\,\text{s}, Wi = 0.08$).

2. Transfer pre-trained weights as initializations for target regimes characterized by high elasticity ($\lambda = 0.70\,\mathrm{s}, Wi = 1.17$) or strong non-linearity ($\alpha = 0.50$).

3. Fine-tune with halved iteration budgets (25,000 iterations), achieving a 50% reduction in computational duration while preventing high-Wi optimization divergence.

Table 4.3 details the hyperparameters and iteration budgets for both precision tiers and initialization strategies.

Table 4.3: Optimization hyperparameters, precision schedule, and transfer learning settings.

Tier	Optimizer	Precision	Iterations (Scratch / TL)	Learning Rate / Line Search
Tier 1	Adam	FP32	40,000 / 20,000	Cosine Annealing ($2.5 \times 10^{-3} \to 2.5 \times 10^{-6}$)
Tier 2	L-BFGS	FP64	10,000 / 5,000	Strong Wolfe ($\alpha_0 = 1.0$, $m = 50$)

4.5 Direct Results

We evaluate the PINN framework on the Oldroyd-B benchmark ($Wi = 0.0833, Re = 0.0417, \lambda = 0.05\,\mathrm{s}, \mu_p = 0.90\,\mathrm{Pa}\cdot\mathrm{s}, \mu_s = 0.10\,\mathrm{Pa}\cdot\mathrm{s}$) under known rheology. Planar velocity observations supervise the stream function $\hat{\psi}$, while the network autonomously reconstructs the unmeasured latent extra-stress tensor $\hat{\boldsymbol{\tau}}(x, y)$ and hydrodynamic pressure $\hat{p}(x, y)$ strictly through the momentum and constitutive residuals.

Crucially, the numerical budget allocated to this direct benchmark differs substantially from the compact schedule adopted for inverse discovery. While inverse parameter identification is standardized on a compact budget of 50,000 iterations (40,000 Adam + 10,000 L-BFGS in Phase 1) to evaluate computational agility and identifiability, the direct benchmark serves to probe the asymptotic accuracy limit of the multi-head architecture across all physical fields on the fine

125k grid. Consequently, the direct benchmark employs an extended two-stage optimization schedule:

1. **Phase 1 (Viscoelastic Stress)**: Trained for 150,000 Adam iterations (FP32) followed by Phase 1 L-BFGS quasi-Newton refinement ($\approx$ 161,500 cumulative iterations), with pressure frozen to cleanly resolve the intense stress boundary layers in the roller gaps.

2. **Phase 2 (Hydrodynamic Pressure)**: Trained for an additional 100,000 Adam iterations with momentum residuals active ($w_{\text{mom}} = 1.0$) and mobile stream function ψ, concluded by 10,250 L-BFGS iterations (FP64) for high-precision pressure gradient reconstruction.

This cumulative investment of over 260,000 optimization iterations produces the reference direct checkpoint (`checkpoint_direct_best_accuracy_phase2`), whose field accuracy metrics are reported in Table 4.4.

4.5.1 Velocity and Stress

Field accuracy is quantified using the relative global L_2 error norm evaluated over all N nodes:

$$\epsilon_{L_2}(\phi) = \frac{\|\hat{\phi} - \phi^{\text{ref}}\|_2}{\|\phi^{\text{ref}}\|_2} = \frac{\sqrt{\sum_{i=1}^{N} |\hat{\phi}(\boldsymbol{x}_i) - \phi^{\text{ref}}(\boldsymbol{x}_i)|^2}}{\sqrt{\sum_{i=1}^{N} |\phi^{\text{ref}}(\boldsymbol{x}_i)|^2}}. \tag{4.19}$$

Complementing the integral L_2 norm, local spatial error distributions in contour maps are reported via the range-relative error:

$$\epsilon_{\text{range}}(\boldsymbol{x}) = \frac{|\hat{\phi}(\boldsymbol{x}) - \phi^{\text{ref}}(\boldsymbol{x})|}{\max_{\Omega}(\phi^{\text{ref}}) - \min_{\Omega}(\phi^{\text{ref}})} \times 100\%, \tag{4.20}$$

which normalizes point-wise discrepancies against the full dynamic span of the field.

Table 4.4 summarizes the validation metrics on the M125k grid.

Table 4.4: Global hydrodynamic field accuracy for the direct Oldroyd-B benchmark ($Wi = 0.0833, Re = 0.0417$) on the M125k evaluation grid (extended schedule: 150k Adam + L-BFGS Phase 1, followed by 100k Adam + 10k L-BFGS Phase 2).

Field Variable	**Symbol**	**Relative L_2 Error**	**Physical Assessment**
Horizontal Velocity	u	**0.98**%	Outflow jets and four-fold symmetry resolved
Vertical Velocity	v	**0.93**%	Exact incompressibility satisfied via ψ
Shear Extra-Stress	τ_{xy}	**1.21**%	Localized peak shear in roller nips resolved
Axial Normal Stress	τ_{xx}	**1.80**%	Elongational stretching along extension axis
Transverse Normal Stress	τ_{yy}	**1.65**%	Compressive stress along inflow streamlines
Hydrodynamic Pressure	p (1-point anchor)	**18.31**%	Anchored at $\boldsymbol{x}_0$; captures global dynamic topology
Gauge-calibrated	p_{cal}	*8.13%*	Mean-adjusted to isolate physical gradient ∇p

Figure 4.5 compares the predicted dimensionless velocity components ($u^* = u/U_{\text{ref}}, v^* = v/U_{\text{ref}}$) against the COMSOL reference (L_2 errors of 0.98% for u and 0.93% for v, with local range-relative errors strictly bounded below 3.5%). The stream-function parameterization maintains the four-fold flow symmetry and accurately locates the central hyperbolic stagnation point.

Figure 4.6 shows the corresponding dimensionless extra-stress components ($\boldsymbol{\tau}^* = \boldsymbol{\tau}/\tau_{\text{ref}}$), with relative L_2 errors of 1.21% on shear stress τ^*_{xy} and 1.80% on normal stress τ^*_{xx}. Stresses are highest in the narrow roller gaps, where intense velocity shearing occurs, and the network reproduces these localized peaks without visible numerical oscillations.

4.5.2 Pressure Recovery

In an enclosed cavity with prescribed boundary velocities, the momentum equations determine the pressure field only through its spatial gradient ∇p. The absolute hydrostatic level p is determined only up to an arbitrary additive constant. To eliminate this freedom during training, the network employs the single-point algebraic hard anchor defined in Eq. (4.15), prescribing the reference Dirichlet value $p(\boldsymbol{x}_0) = p_{\text{ref}}$ at a single external wall node $\boldsymbol{x}_0$.

From a hydrodynamic and rheological perspective, therefore, the fundamen-

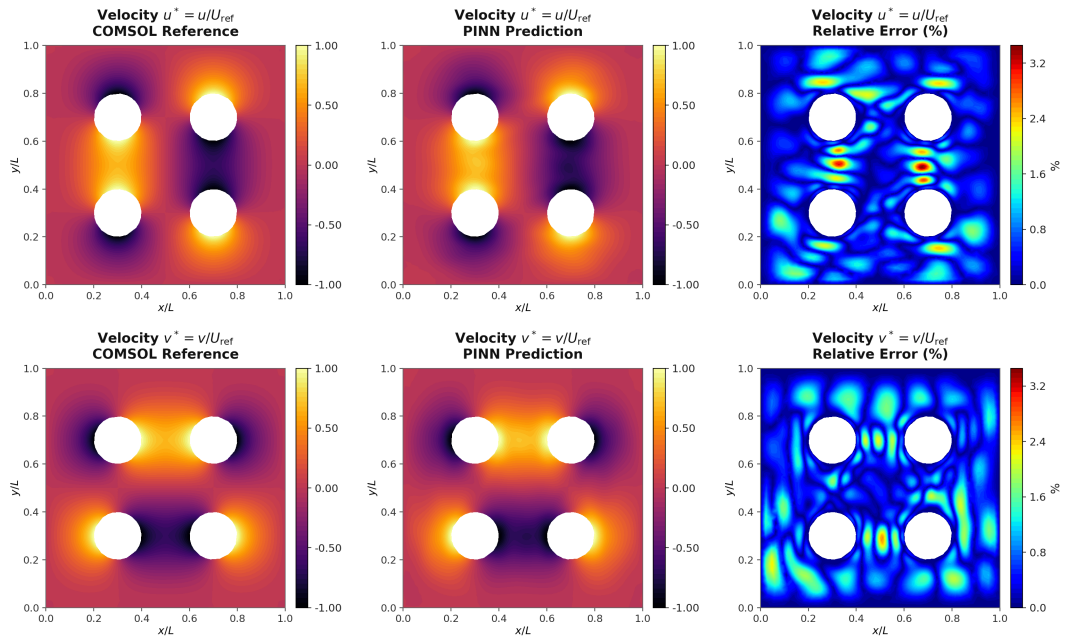


Figure 4.5: Direct problem kinematic validation for the canonical Oldroyd-B benchmark ($\lambda = 0.05\,\mathrm{s}$, $\mu_p = 0.90\,\mathrm{Pa}\cdot\mathrm{s}$, $\mu_s = 0.10\,\mathrm{Pa}\cdot\mathrm{s}$, $Wi = 0.0833$, $Re = 0.0417$, $\alpha = 0$, $\varepsilon = 0$): 3-column comparison of dimensionless horizontal ($u^* = u/U_{\mathrm{ref}}$) and vertical ($v^* = v/U_{\mathrm{ref}}$) velocity fields on normalized coordinates $[0, 1]^2$ across COMSOL reference, PINN prediction, and range-relative error on the M125k mesh.

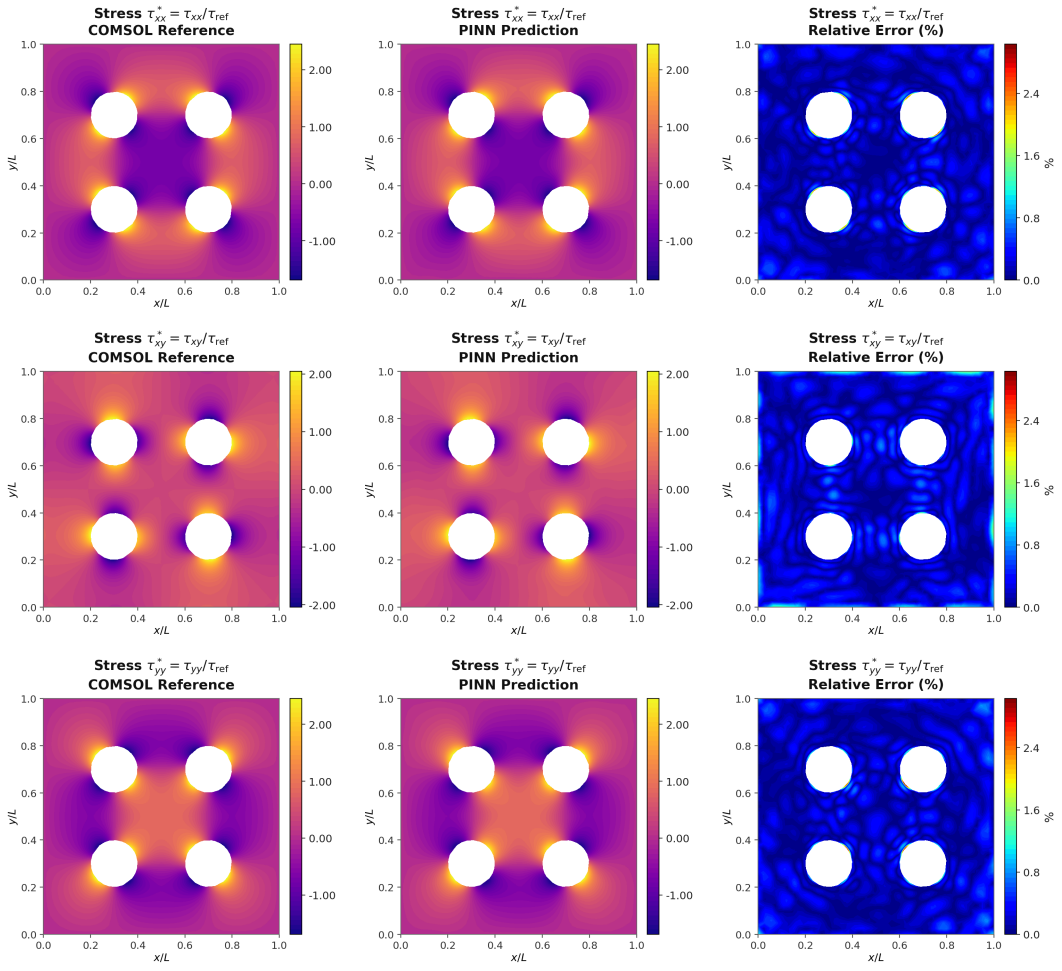


Figure 4.6: Direct problem viscoelastic extra-stress tensor validation for the canonical Oldroyd-B benchmark ($\lambda = 0.05\,\mathrm{s}$, $\mu_p = 0.90\,\mathrm{Pa}\cdot\mathrm{s}$, $\mu_s = 0.10\,\mathrm{Pa}\cdot\mathrm{s}$, $Wi = 0.0833$, $Re = 0.0417$, $\alpha = 0$, $\varepsilon = 0$): 3-column comparison of dimensionless COMSOL ground truth ($\boldsymbol{\tau}^* = \boldsymbol{\tau}/\tau_{\mathrm{ref}}$), PINN predictions, and range-relative errors for normal components (τ^*_{xx}, τ^*_{yy}) and shear component τ^*_{xy} on normalized coordinates $[0,1]^2$.

tal criterion for solver correctness is whether the physics-informed architecture accurately learns the true spatial pressure gradient ∇p, rather than matching an arbitrary hydrostatic reference constant, also given that the other fields are completely indepentend from pressure shifts.

Figure 4.7 displays the dimensionless hydrodynamic pressure field $p^*(\boldsymbol{x}) = p(\boldsymbol{x})/p_{\text{ref}}$ predicted directly by the network under this single-point anchor. Over the entire 125k evaluation mesh, the raw anchored field achieves a relative L_2 error of 18.31%.

This 18.31% discrepancy represents the operational performance of the solver in a realistic experimental setting where only a single physical pressure tap is available at $\boldsymbol{x}_0$. Because the global two-dimensional pressure surface is pinned at an isolated peripheral boundary node, any infinitesimal residual discrepancy in the learned local derivative ∇p at $\boldsymbol{x}_0$ shifts the entire 2D surface vertically by a uniform constant offset across the cavity. As documented in Table 4.4, removing this uniform hydrostatic offset via domain-mean calibration ($p_{\text{cal}} = p - \bar{p} + \bar{p}_{\text{ref}}$) reduces the relative L_2 error from 18.31% to 8.13%, while the mean range-relative error across the cavity is 4.08% (Figure 4.8).

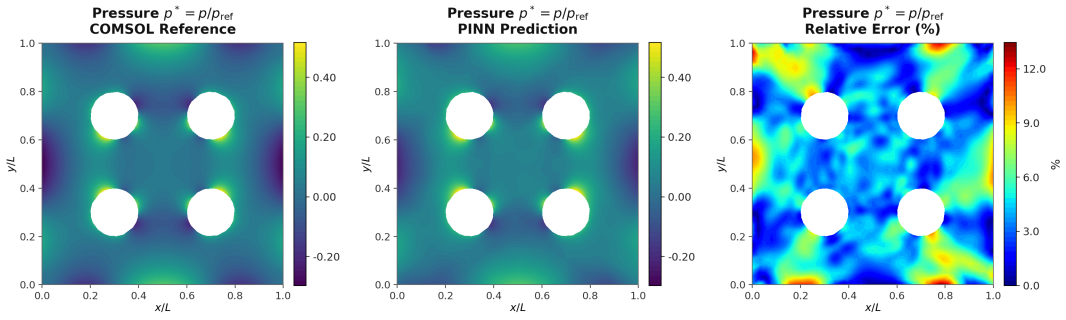


Figure 4.7: Official hydrodynamic pressure field reconstruction in Direct Mode under single-point algebraic anchoring ($p(\boldsymbol{x}_0) = p_{\text{ref}}$) for the canonical Oldroyd-B benchmark ($\lambda = 0.05\,\text{s}$, $\mu_p = 0.90\,\text{Pa}\cdot\text{s}$, $\mu_s = 0.10\,\text{Pa}\cdot\text{s}$, $Wi = 0.0833$, $Re = 0.0417$): 3-column comparison of dimensionless COMSOL ground truth ($p^* = p/p_{\text{ref}}$), raw PINN prediction, and range-relative error distributions on normalized coordinates $[0, 1]^2$ (L_2 error of 18.31%, mean range-relative error of 4.08%).

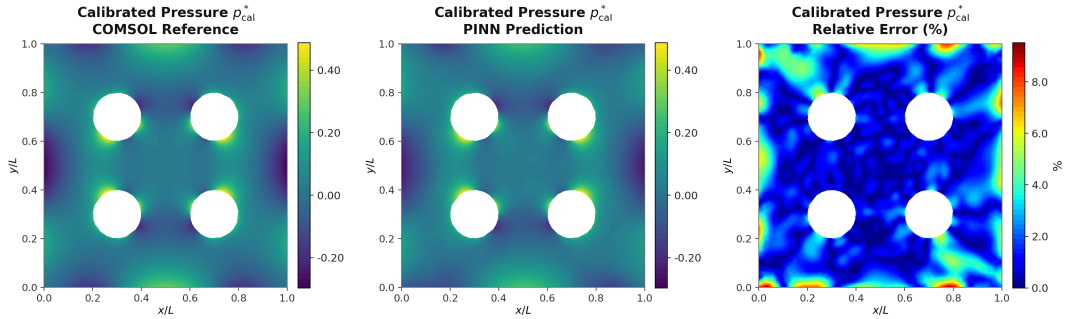


Figure 4.8: Hydrodynamic pressure field after gauge calibration ($p_{\text{cal}} = p - \bar{p} + \bar{p}_{\text{ref}}$) for the Oldroyd-B benchmark ($\lambda = 0.05\,\text{s}$, $\mu_p = 0.90\,\text{Pa}\cdot\text{s}$, $\mu_s = 0.10\,\text{Pa}\cdot\text{s}$, $Wi = 0.0833$, $Re = 0.0417$): 3-column validation isolating the physical gradient ∇p (L_2 error reduced to 8.13%, mean range-relative error of 4.08%). The PINN faithfully captures the four-fold symmetric suction troughs in the roller nips and the central stagnation saddle.

The comparison between the raw anchored prediction (Figure 4.7) and the gauge-calibrated distribution (Figure 4.8) illustrates the role of the shift constant: all the relevant physical features are reproduced consistently. The error is concentrated in a constant vertical offset induced by the single-point boundary anchor, showing that the network recovers the physical pressure gradient ∇p governing the momentum balance.

4.6 Inverse Parameter Discovery

In inverse discovery mode, the PINN operates as a virtual rheometer: no internal stress measurements are supplied. The network assimilates domain velocity data, solid boundary kinematic conditions, and the roller stress anchor ($\mathcal{L}_{\text{roll_stress}}$). Internal stresses and pressure remain unobserved, inferred simultaneously alongside unknown fluid parameters.

4.6.1 Oldroyd-B and Giesekus Benchmarks

We evaluate four canonical benchmarks spanning linear Hookean and non-linear shear-thinning response:

- **Benchmark 1 (Oldroyd-B from Scratch)**: Modest elasticity ($Wi = 0.08$), trained ex-novo for 50,000 iterations starting from blind offset guesses ($\lambda_0 = 0.0535\,\mathrm{s}$ and $\mu_p^0 = 0.963\,\mathrm{Pa}\cdot\mathrm{s}$, both +7% from truth).

- **Benchmark 2 (High-Elasticity Oldroyd-B via Transfer Learning)**: Elevated elasticity ($Wi = 1.17, \lambda = 0.70\,\mathrm{s}$), fine-tuned via TL for 25,000 iterations.

- **Benchmark 3 (Non-linear Giesekus with Blind Guess from Scratch)**: Non-linear mobility ($\alpha^{\mathrm{true}} = 0.35$), initialized from a blind guess $\alpha_0 = 0.25$ and trained for 50,000 iterations.

- **Benchmark 4 (Extreme-Mobility Giesekus via Transfer Learning)**: Maximum physical mobility ($\alpha^{\mathrm{true}} = 0.50, \lambda = 0.30\,\mathrm{s}, \mu_p = 0.05\,\mathrm{Pa}\cdot\mathrm{s}$), fine-tuned via TL for 25,000 iterations.

Table 4.5 summarizes the discovery accuracy and kinematic errors across these benchmarks.

Table 4.5: Inverse parameter discovery results across the four representative benchmarks on the M125k grid.

#	Model	Strategy	Target Parameter	PINN Estimation	Error	$L_2(u,v)$
1	**Oldroyd-B**	Ex-Novo	$\alpha = 0.000$	**0.0003**	$\approx$ **0**	
	($Wi = 0.08$)	(50k iters)	$\lambda = 0.050\,\mathrm{s}$	**0.0502** s	+**0.41**%	0.04%
			$\mu_p = 0.900\,\mathrm{Pa}\cdot\mathrm{s}$	**0.9049** Pa · s	+**0.54**%	
2	**Oldroyd-B**	Transfer Learning	$\alpha = 0.000$	**0.0011**	$\approx$ **0**	
	($Wi = 1.17$)	(25k iters, −50%)	$\lambda = 0.700\,\mathrm{s}$	**0.7089** s	+**1.27**%	0.33%
			$\mu_p = 0.500\,\mathrm{Pa}\cdot\mathrm{s}$	**0.5019** Pa · s	+**0.38**%	
3	**Giesekus**	Ex-Novo (Blind)	$\alpha = 0.350$	**0.3298**	−**5.76**%	
	($\alpha_0 = 0.25$)	(50k iters)	$\lambda = 0.100\,\mathrm{s}$	**0.1032** s	+**3.16**%	0.46%
			$\mu_p = 0.500\,\mathrm{Pa}\cdot\mathrm{s}$	**0.5307** Pa · s	+**6.14**%	
4	**Giesekus**	Transfer Learning	$\alpha = 0.500$	**0.4665**	−**6.70**%	
		(25k iters, −50%)	$\lambda = 0.300\,\mathrm{s}$	**0.3077** s	+**2.57**%	0.52%
			$\mu_p = 0.050\,\mathrm{Pa}\cdot\mathrm{s}$	**0.0522** Pa · s	+**4.40**%	

Parameter estimation errors for Oldroyd-B remain below 1.3% across both low and high Weissenberg numbers ($Wi = 0.08$ and $Wi = 1.17$). In both Oldroyd-B benchmarks, the mobility parameter converges to zero ($\alpha_{\text{est}} = 0.0003$ and 0.0011, respectively), demonstrating that the network correctly identifies the absence of shear-thinning without prior assumptions. For Giesekus, the solver simultaneously inverts three coupled parameters (α, λ, μ_p), identifying mobility α within 6.7% error even from a blind initialization ($\alpha_0 = 0.25$ vs 0.35).

Inverse Stress Fields

To verify that the inverse PINN reconstructs continuous physical fields rather than fitting isolated scalar parameters, Figures 4.9 and 4.10 present two-dimensional contour comparisons of the unobserved latent extra-stress components $(\tau^*_{xx}, \tau^*_{xy}, \tau^*_{yy})$ for Benchmark 1 (Oldroyd-B, $Wi = 0.08$) and Benchmark 3 (Giesekus, $\alpha = 0.35$).

As illustrated in Figure 4.9 for the Oldroyd-B benchmark, the inverse solver reconstructs the unmeasured stress fields across the domain. Normal stresses achieve relative L_2 errors of 0.64% on τ^*_{xx} and 0.70% on τ^*_{yy}, with domain-wide mean range-relative errors below 0.15% (maximum localized range error below 2.66%). Shear stress τ^*_{xy} is recovered with a mean range-relative error of 4.35%, concentrated along high-shear boundary interfaces where velocity gradients are steepest.

Figure 4.10 evaluates the more challenging non-linear Giesekus inverse problem (Benchmark 3), where the solver simultaneously inverts three strongly coupled parameters (α, λ, μ_p) starting from an offset guess ($\alpha_0 = 0.25$ vs. $\alpha^{\text{true}} = 0.35$). Even under this non-linear constitutive coupling, the reconstructed stress components remain consistent with the reference solution: relative L_2 errors evaluate to 7.28% for τ^*_{xx}, 5.92% for shear stress τ^*_{xy}, and 6.85% for transverse normal stress τ^*_{yy}. Over the entire domain, the mean range-relative errors remain

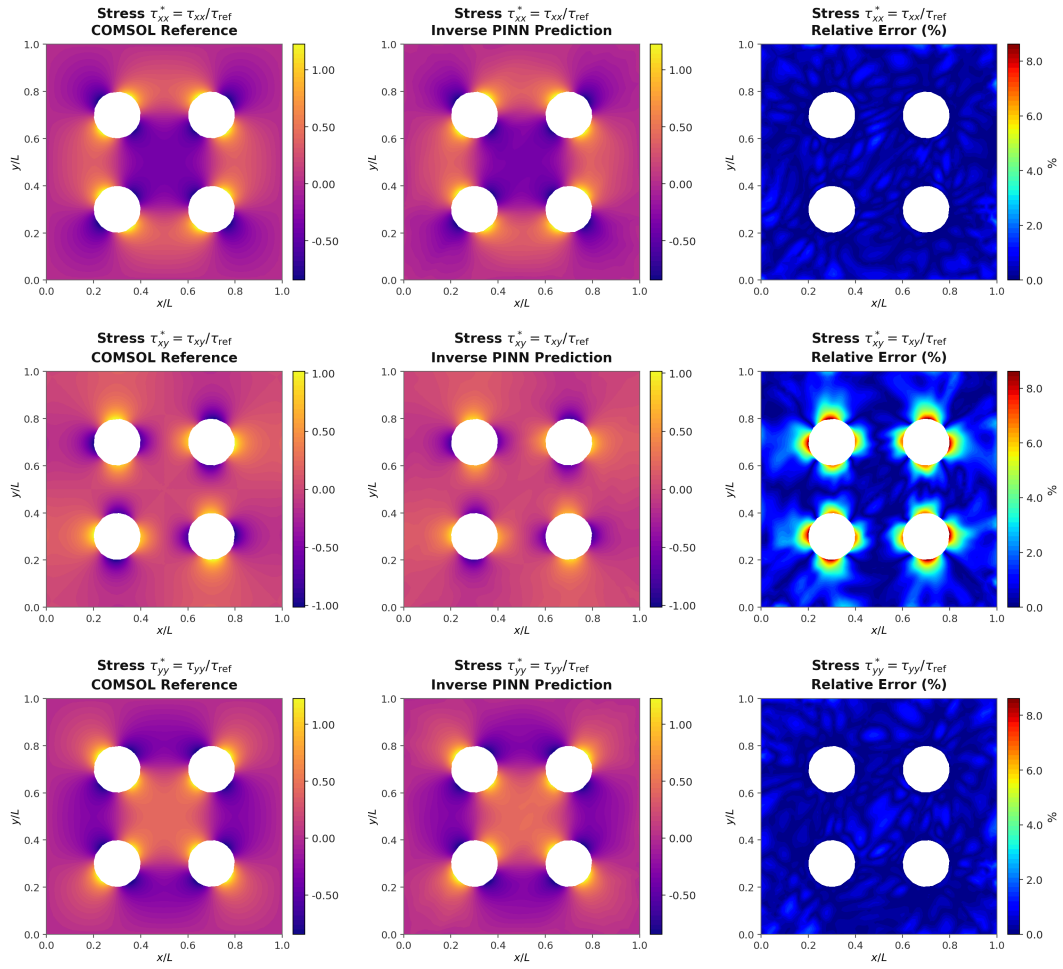


Figure 4.9: Inverse Mode latent viscoelastic extra-stress tensor reconstruction for the canonical Oldroyd-B fluid (Benchmark 1: $\lambda^{\text{true}} = 0.050\,\text{s}$, $\mu_p^{\text{true}} = 0.900\,\text{Pa}\cdot\text{s}$, $\mu_s^{\text{true}} = 0.100\,\text{Pa}\cdot\text{s}$, $Wi = 0.0833$, $Re = 0.0417$, $\alpha^{\text{true}} = 0$, $\varepsilon^{\text{true}} = 0$): 3-column comparison of dimensionless COMSOL reference ($\boldsymbol{\tau}^* = \boldsymbol{\tau}/\tau_{\text{ref}}$), Inverse PINN predictions, and range-relative errors across normal components (τ^*_{xx}, τ^*_{yy}) and shear component τ^*_{xy} on normalized coordinates $[0,1]^2$. Note the high spatial precision on normal stresses (L_2 error of 0.64% on τ^*_{xx} and 0.70% on τ^*_{yy}, with mean range errors below 0.15%).

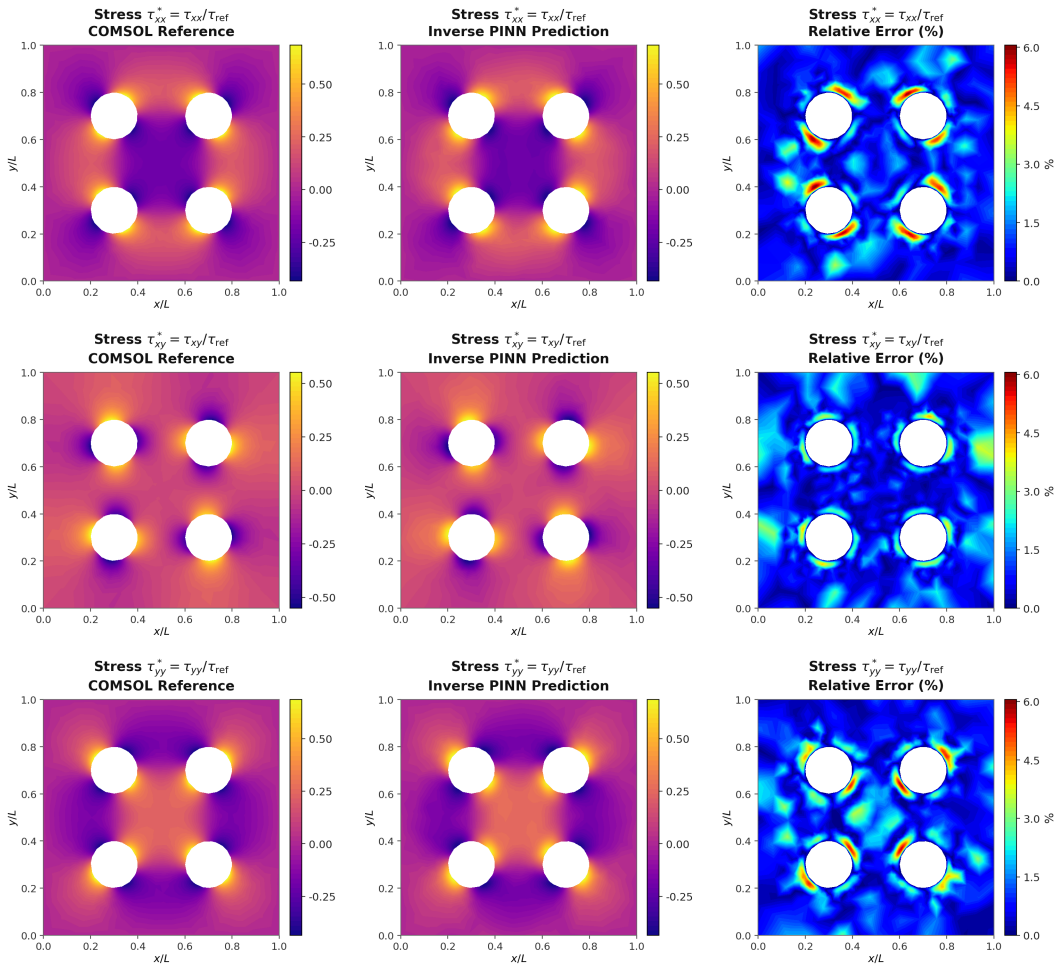


Figure 4.10: Inverse Mode latent viscoelastic extra-stress tensor reconstruction for the non-linear shear-thinning Giesekus fluid (Benchmark 3: $\alpha^{\text{true}} = 0.350$, $\lambda^{\text{true}} = 0.100\,\text{s}$, $\mu_p^{\text{true}} = 0.500\,\text{Pa}\cdot\text{s}$, $\mu_s^{\text{true}} = 0.500\,\text{Pa}\cdot\text{s}$, $Wi = 0.1667$, $Re = 0.0417$, $\varepsilon^{\text{true}} = 0$, blind initial guess $\alpha_0 = 0.25$): 3-column comparison across COMSOL ground truth, Inverse PINN predictions, and range-relative errors for normal components $(\tau^*_{xx}, \tau^*_{yy})$ and shear component τ^*_{xy}. Relative L_2 errors remain bounded between 5.92% and 7.28%, with mean range-relative errors strictly below 1.55%, confirming simultaneous parameter discovery and stress field recovery.

strictly below 1.55% across all three stress components. The solver successfully reproduces the non-linear saturation plateau near the central stagnation point and resolves the shear-thinning stress redistribution around the rotating cylinders without introducing spurious oscillations or unphysical stress accumulation. These results confirm that the inverse PINN resolves both boundary-layer shear and extensional stagnation dynamics simultaneously with parameter identification.

Figure 4.11 illustrates parameter trajectories and objective function evolution during staged Adam (FP32, initial 40k epochs) and L-BFGS (FP64, final 10k iterations) optimization on the M125k dataset. In Figure 4.11(left), the normalized parameter trajectories $\lambda/\lambda_{\mathrm{GT}}$ and $\mu_p/\mu_{p,\mathrm{GT}}$ initially rise from 1.07 to an overshoot peak near 1.28 during the first 2,500 Adam iterations before reversing course and converging monotonically toward 1.00. This transient behavior has a direct physical cause: at iteration zero, the stress network $\mathcal{N}_{\boldsymbol{\tau}}$ is initialized near zero ($\boldsymbol{\tau} \approx \mathbf{0}$). In the absence of a structured stress field, the constitutive residual gradients momentarily push parameter estimates upward to balance the prescribed deformation rate. Once the stress head establishes the localized boundary layers and shear bands (around epoch 3,000), the optimizer reorients toward the physical minimum, guiding parameters into the physical neighborhood. Upon switching to the second-order L-BFGS optimizer at iteration 40k, the kinematic data loss drops sharply by an order of magnitude (from 10^{-6} to 1.81×10^{-7}), while the constitutive loss settles monotonically to $\mathcal{L}_{\mathrm{con}} = 6.13 \times 10^{-6}$, providing smooth asymptotic settling without numerical oscillations.

4.6.2 Transfer Learning Performance

The benchmarks in Table 4.5 demonstrate two practical benefits of transfer learning:

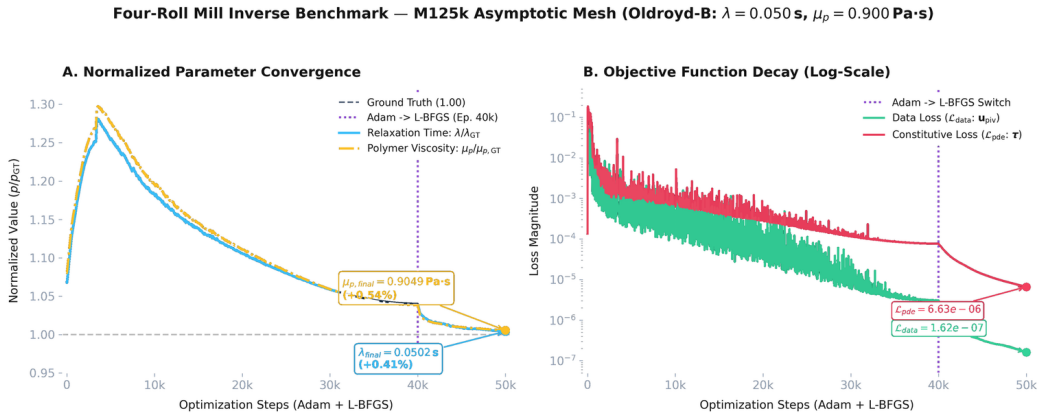


Figure 4.11: Parameter convergence trajectories (λ and μ_p, left) and staged objective loss decay ($\mathcal{L}_{\text{data}}$ and $\mathcal{L}_{\text{con}}$, right) during Adam (FP32) and L-BFGS (FP64) optimization on the M125k dataset for the canonical Oldroyd-B benchmark ($\lambda^{\text{true}} = 0.050\,\text{s}$, $\mu_p^{\text{true}} = 0.900\,\text{Pa}\cdot\text{s}$, $\mu_s^{\text{true}} = 0.100\,\text{Pa}\cdot\text{s}$, $Wi = 0.0833$, $Re = 0.0417$).

1. **Computational Acceleration**: Warm-starting from a pre-trained baseline reduces required iterations by half (25,000 vs. 50,000), halving training time.

2. **Optimization Regularization**: Initializing weights from a lower-Wi solution places the optimization in a physically reasonable neighborhood, avoiding the poor local minima that can occur when training from scratch at higher elasticity ($Wi = 1.17$) or high mobility ($\alpha = 0.50$).

4.6.3 Collocation Density Sensitivity

To examine how collocation sampling density affects PINN parameter discovery, we conducted a systematic grid study across five resolutions from 5k to 125k nodes (Table 4.6 and Figure 4.12) on an Oldroyd-B case ($\lambda = 0.100\,\text{s}, \mu_p = 0.500\,\text{Pa}\cdot\text{s}$).

Across a 25-fold variation in collocation point density (from 5,086 to 125,456 nodes), parameter estimation errors remain consistently bounded within a narrow interval (∼0.8%–8.8%), while velocity fields are reconstructed with high precision ($L_2 \leq 0.58\%$).

Table 4.6: Collocation density sensitivity and parameter estimation accuracy across five resolutions ($\lambda^{\text{true}} = 0.100\,\text{s}$, $\mu_p^{\text{true}} = 0.500\,\text{Pa}\cdot\text{s}$).

Mesh Resolution	Evaluation Nodes	λ_{est} [s]	λ Error	$\mu_{p,\text{est}}$ [Pa·s]	μ_p Error	Velocity L_2
5k (Coarse)	**5,086**	**0.1008**	**+0.79%**	**0.5103**	**+2.06%**	**0.15%**
12k	12,760	0.1071	+7.10%	0.5389	+7.78%	0.55%
29k	29,392	0.1088	+8.79%	0.5455	+9.10%	0.58%
52k	52,648	0.1062	+6.18%	0.5323	+6.46%	0.44%
125k (Dense)	125,456	0.1083	+8.32%	0.5412	+8.24%	0.50%

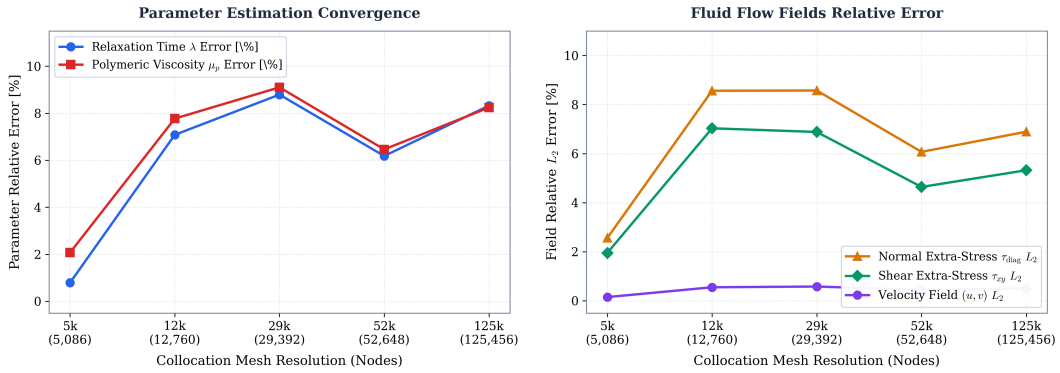


Figure 4.12: Collocation density study: comparison of parameter estimation errors (λ and μ_p) and velocity L_2 errors across resolutions from 5,086 to 125,456 collocation points for the Oldroyd-B benchmark ($\lambda^{\text{true}} = 0.100\,\text{s}$, $\mu_p^{\text{true}} = 0.500\,\text{Pa}\cdot\text{s}$, $\mu_s^{\text{true}} = 0.500\,\text{Pa}\cdot\text{s}$, $Wi = 0.1667$, $Re = 0.0417$, $\alpha = 0$, $\varepsilon = 0$).

The modest variations observed across grid resolutions reflect minor random fluctuations associated with stochastic optimization and spatial collocation sampling rather than systematic grid-dependent bias. All resolutions yield comparable, reliable parameter estimates, demonstrating that the inverse solver is robust to collocation point density and does not require excessively dense grids to achieve satisfactory convergence.

4.7 Physical Limits and Identifiability Boundaries

4.7.1 Creeping Flow Hydrodynamic Invariance

In Phase 2, estimating solvent viscosity μ_s alongside pressure p by minimizing momentum residuals reveals an identifiability boundary. In our flow conditions, the Reynolds number is small ($Re = 0.0417 \ll 1$), placing the flow in the creeping Stokes regime where convective inertia is negligible:

$$-\nabla p + \mu_s \nabla^2 \boldsymbol{u} + \nabla \cdot \boldsymbol{\tau} = \mathbf{0}. \tag{4.21}$$

Variations in μ_s do not alter the kinematic velocity field; they simply shift the pressure gradient ∇p required to balance the viscous forces.

Figure 4.13 shows the velocity profiles along the three cutlines (Horizontal, Vertical, and Diagonal) computed from COMSOL on the M125k mesh at $\lambda = 0.05\,\mathrm{s}$ for three solvent viscosities: $\eta_s = 0.10$, 0.30, and $0.50\,\mathrm{Pa}\cdot\mathrm{s}$ (viscosity ratios $\beta = 0.10, 0.30$, and 0.50). The velocity profiles overlap across all three cutlines, with local differences below 0.03%. The kinematic field is therefore independent of η_s, which leaves the loss gradient $\partial \mathcal{L}_{\mathrm{mom}}/\partial \mu_s$ uninformative and prevents gradient-based optimization of μ_s from velocity measurements alone

without independent pressure or traction data.

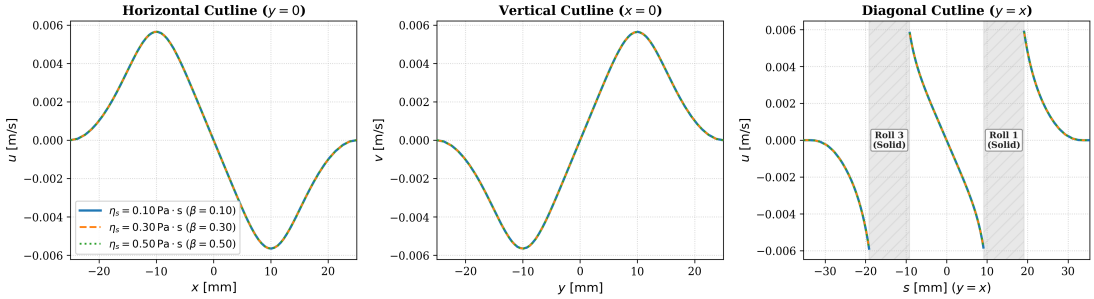


Figure 4.13: Hydrodynamic invariance in creeping flow: velocity profiles along the Horizontal ($y = 0$), Vertical ($x = 0$), and Diagonal ($y = x$) cutlines for three solvent viscosities ($\eta_s = 0.10, 0.30, 0.50\,\mathrm{Pa \cdot s}$) in an Oldroyd-B fluid on the M125k mesh ($\lambda = 0.05\,\mathrm{s}$, $\mu_p = 0.90$–$0.50\,\mathrm{Pa \cdot s}$, $Wi = 0.0833$, $Re = 0.0417$). The velocity profiles overlap, confirming kinematic independence from η_s under prescribed cylinder velocities.

4.7.2 Identifiability of Linear PTT

When applying the inverse PINN solver to synthetic data generated from the Phan-Thien–Tanner (PTT) model, parameter degradation occurs: the extensibility parameter collapses toward zero ($\varepsilon_{\mathrm{est}} \approx 10^{-4}$), while relaxation time λ and polymeric viscosity μ_p are underestimated by approximately 58–61% (Table 4.7). Despite these individual errors, the reconstructed velocity field matches COMSOL closely ($L_2 < 0.9\%$).

Table 4.7: PTT inverse problem parameter degradation and kinematic error under boundary velocity supervision.

Benchmark Scenario	$\varepsilon^{\mathrm{true}} \to \varepsilon_{\mathrm{est}}$ **(Error)**	$\lambda^{\mathrm{true}} \to \lambda_{\mathrm{est}}$ **(Error)**	$\mu_p^{\mathrm{true}} \to \mu_{p,\mathrm{est}}$ **(Error)**	**Velocity** L_2
Linear PTT ($\lambda = 1.00\,\mathrm{s}, \varepsilon = 0.30$) ($\mu_p = 0.500\,\mathrm{Pa \cdot s}, Wi = 1.67$)	$0.300 \to 1.7 \times 10^{-4}$ (**−99.94**%)	$1.000 \to 0.421\,\mathrm{s}$ (**−57.89**%)	$0.500 \to 0.213\,\mathrm{Pa \cdot s}$ (**−57.44**%)	**0.79**%
Linear PTT ($\lambda = 0.70\,\mathrm{s}, \varepsilon = 0.50$) ($\mu_p = 0.950\,\mathrm{Pa \cdot s}, Wi = 1.17$)	$0.500 \to 4.4 \times 10^{-4}$ (**−99.91**%)	$0.700 \to 0.270\,\mathrm{s}$ (**−61.43**%)	$0.950 \to 0.369\,\mathrm{Pa \cdot s}$ (**−61.19**%)	**0.88**%

This behavior highlights an open question regarding the identifiability of the linear PTT formulation under planar kinematic observations. In the viscoelastic PINN literature implementations (such as ViscoelasticNet) have successfully

identified PTT constitutive parameters, confirming that the differential operator itself does not inherently preclude convergence in physics-informed frameworks.

Consequently, the lack of parameter convergence observed in the present two-dimensional four-roll configuration requires further investigation. Rather than indicating an absolute physical barrier, this outcome suggests that the loss landscape is highly sensitive to the chosen parameter ranges, non-dimensional scales, and relative loss weighting. The exploration of alternative convergence strategies and optimization regimes remains an active direction of research to determine the conditions under which the PTT parameters can be reliably resolved in this flow geometry.

Chapter 5

Conclusions and Future Perspectives

This thesis evaluated a Physics-Informed Neural Network (PINN) framework for the modeling and inverse rheological characterization of steady, two-dimensional viscoelastic flows in a four-roll mill geometry. By imposing conservation of mass, momentum balance, and constitutive laws into a decoupled multi-head architecture, we investigated field reconstruction and inverse parameter estimation across three constitutive models with increasing nonlinearity: Oldroyd-B, Giesekus, and linear Phan-Thien-Tanner (PTT) fluids.

To capture velocity, pressure, and stress in a closed recirculating cavity, we used a three-head architecture ($\mathcal{N}_\psi, \mathcal{N}_p, \mathcal{N}_{\boldsymbol{\tau}}$) with SiLU activations and zero output initialization for stress and pressure. Formulating the velocity through a stream function $\psi(x, y)$ enforces incompressibility ($\nabla \cdot \boldsymbol{u} \equiv 0$) by construction, while pressure gauge indeterminacy is fixed by an algebraic single-point anchor. Training is divided into two stages: Phase 1 optimizes kinematics and viscoelastic stresses, while Phase 2 optimizes momentum balance and pressure while keeping the stream function trainable. Combining FP32 Adam for initial exploration with FP64 L-BFGS for higher-order derivatives provided stable convergence without

optimization plateaus.

In direct mode on the canonical Oldroyd-B benchmark ($Wi = 0.0833, Re = 0.0417$), velocity supervision alone allowed the network to reconstruct extra-stresses and pressure across the domain. Compared to the COMSOL reference solution on 125,456 nodes, relative L_2 errors were 0.98% for u and 0.93% for v. Extra-stress errors were 1.21% for τ_{xy}, 1.80% for τ_{xx}, and 1.65% for τ_{yy}, resolving localized stress peaks in roller gaps. The pressure field matched reference profiles with an 8.13% range-relative error after gauge calibration.

In inverse mode, the network identified fluid parameters using planar velocity data and roller boundary stress anchoring, without interior stress measurements. For Oldroyd-B fluids, relaxation time λ and polymeric viscosity μ_p were identified within 0.41%–1.27% error across low and moderate elasticity regimes ($Wi = 0.08$ and 1.17). For shear-thinning Giesekus fluids, the solver estimated three parameters (α, λ, μ_p) simultaneously, recovering mobility α within 5.76%–6.70% error even up to the theoretical limit $\alpha = 0.50$. Transfer learning cut training iterations from 50,000 to 25,000 steps and stabilized convergence at higher Weissenberg numbers, while parameter estimates remained accurate down to 5,086 collocation points.

The analysis also identifies several operational limitations. Breaking the scaling ambiguity between polymeric viscosity μ_p and stress magnitude requires stress data on the cylinder surfaces. While this avoids domain-wide stress measurements, measuring roller surface stresses experimentally requires optical birefringence or traction sensors. In addition, when velocity data are concentrated near the rollers, parameters governing extensional saturation, such as the PTT extensibility ε, become weakly observable. While the single-point algebraic anchor mathematically fixes the pressure gauge, local errors in ∇p at the anchor node shift the baseline across the cavity, yielding an 18.31% raw relative error compared

to 8.13% after removing the mean. Finally, the analysis is limited to steady, two-dimensional flows at moderate elasticity ($Wi \leq 1.67$), without accounting for three-dimensional end-wall effects or elastic instabilities at higher deformation rates.

These physical and operational findings highlight several practical extensions for future research, which can be organized into four primary directions:

1. **Changing flow field parameter**: The present benchmark relies on four identical counter-rotating cylinders, which enforce purely extensional flow at the cavity center. Future work should evaluate asymmetric or time-dependent boundary kinematics, such as independent cylinder rotation speeds or alternating drive ratios (varying the kinematic parameter k). Testing these configurations will verify whether the decoupled architecture adapts to mixed shearing, rotational, and and extensional dynamics.

2. **Addition of Gaussian Noise**: The synthetic datasets used throughout this work represent ideal, noise-free velocity fields. In physical Particle Image Velocimetry (PIV), observations are inevitably affected by optical noise and tracer tracking uncertainties. Adding controlled Gaussian noise to planar velocity and roller boundary observations will establish the sensitivity of automatic differentiation under realistic experimental conditions, determining the signal-to-noise threshold over which this framework it no more applicable.

3. **Exponential Constitutive PTT**: The current study implemented linear PTT fluids, in which the stress trace enters linearly as $1 + (\varepsilon\lambda/\mu_p)\mathrm{tr}(\boldsymbol{\tau})$. Extending the formulation to the exponential PTT (ePTT) model, where $Y(\mathrm{tr}(\boldsymbol{\tau})) = \exp[(\varepsilon\lambda/\mu_p)\mathrm{tr}(\boldsymbol{\tau})]$, introduces steeper non-linear gradients during backpropagation. Investigating ePTT will test whether the staged optimization scheme remains stable when confronted with exponential stress

growth near stagnation lines.

4. **Convergence Mapping Across Parameter Space**: Systematic parameter sweeps across Weissenberg numbers (Wi), Reynolds numbers (Re), and objective loss weightings are needed to map the complete convergence range of this framework. This research will identify the precise boundaries between smooth parameter recovery and numerical stiffness, defining the operational window where this physics-informed framework is most reliable.

Physics-informed neural networks offer a practical framework to combine kinematic measurements with constitutive modeling in viscoelastic flows. By embedding conservation laws and objective stress transport into decoupled network branches, the model reconstructs unobserved stress and pressure fields directly from planar velocity data. The ability to use PINNs as a virtual rheometer suggests that high-speed, non-intrusive characterization of polymer melts is possible using only velocity measurements (PIV), potentially reducing the cost of quality control in industrial processes. By resolving localized boundary stresses, reproducing central extensional features, and estimating constitutive parameters without a numerical mesh, the framework provides a foundation for data-driven rheological diagnostics in viscoelastic fluids processing.

Bibliography

[1] R. Byron Bird, Robert C. Armstrong, and Ole Hassager. *Dynamics of Polymeric Liquids, Volume 1: Fluid Mechanics*. John Wiley & Sons, 2nd edition, 1987.

[2] Christopher W. Macosko. *Rheology: Principles, Measurements, and Applications*. Wiley-VCH, 1994.

[3] Robert G. Owens and Timothy N. Phillips. *Computational Rheology*. Imperial College Press, 2002.

[4] Geoffrey I. Taylor. The formation of emulsions in definable fields of flow. *Proceedings of the Royal Society of London. Series A*, 146(858):501–523, 1934.

[5] G. G. Fuller and L. G. Leal. Flow birefringence of dilute polymer solutions in two-dimensional flows. *Journal of Non-Newtonian Fluid Mechanics*, 6(3-4):193–210, 1980.

[6] Gerald G. Fuller and L. Gary Leal. Flow birefringence of concentrated polymer solutions in two-dimensional flows. *Journal of Polymer Science: Polymer Physics Edition*, 19(4):557–587, 1981.

[7] B. J. Bentley and L. G. Leal. An experimental investigation of drop deformation and breakup in steady, two-dimensional linear flows. *Journal of Fluid Mechanics*, 167:241–283, 1986.

[8] Maziar Raissi, Paris Perdikaris, and George E. Karniadakis. Physics-informed neural networks: A deep learning framework for solving forward and inverse problems involving nonlinear partial differential equations. *Journal of Computational Physics*, 378:686–707, 2019.

[9] Salvatore Cuomo, Vincenzo Schiano Di Cola, Fabio Giampaolo, Gianluigi Rozza, Maziar Raissi, and Francesco Piccialli. Scientific machine learning through physics-informed neural networks: Where we are and what's next. *Journal of Scientific Computing*, 92(3):88, 2022.

[10] Atilim Gunes Baydin, Barak A. Pearlmutter, Alexey Andreyevich Radul, and Jeffrey Mark Siskind. Automatic differentiation in machine learning: a survey. *Journal of Machine Learning Research*, 18(153):1–43, 2018.

[11] Sukirt Thakur, Maziar Raissi, and Arezoo M. Ardekani. Viscoelasticnet: A physics informed neural network framework for stress discovery and model selection. *Journal of Non-Newtonian Fluid Mechanics*, 330:105265, 2024.

[12] COMSOL AB. *COMSOL Multiphysics Reference Manual, version 6.0*. COMSOL AB, Stockholm, Sweden, 2021.

[13] James G. Oldroyd. On the formulation of rheological equations of state. *Proceedings of the Royal Society of London. Series A*, 200(1063):523–541, 1950.

[14] Hanswalter Giesekus. A simple constitutive equation for polymer fluids based on the concept of deformation-dependent tensorial mobility. *Journal of Non-Newtonian Fluid Mechanics*, 11(1-2):69–109, 1982.

[15] Nhan Phan-Thien and Roger I. Tanner. A new constitutive equation derived from network theory. *Journal of Non-Newtonian Fluid Mechanics*, 2(4):353–365, 1977.

[16] Gaetano D'Avino, Martien A. Hulsen, and Pier Luca Maffettone. Numerical simulations of dispersive mixing of viscoelastic suspensions in a four-roll mill. *Journal of Non-Newtonian Fluid Mechanics*, 247:176–187, 2017.

[17] Ian Goodfellow, Yoshua Bengio, and Aaron Courville. *Deep Learning*. MIT Press, 2016.

[18] George Cybenko. Approximation by superpositions of a sigmoidal function. *Mathematics of Control, Signals and Systems*, 2(4):303–314, 1989.

[19] Kurt Hornik, Maxwell Stinchcombe, and Halbert White. Multilayer feedforward networks are universal approximators. *Neural Networks*, 2(5):359–366, 1989.

[20] Diederik P. Kingma and Jimmy Ba. Adam: A method for stochastic optimization. *arXiv preprint arXiv:1412.6980*, 2014.

[21] Jorge Nocedal. Updating quasi-newton matrices with limited storage. *Mathematics of Computation*, 35(151):773–782, 1980.

[22] Dong C. Liu and Jorge Nocedal. On the limited memory bfgs method for large scale optimization. *Mathematical Programming*, 45(1-3):503–528, 1989.

[23] Sifan Wang, Yujun Teng, and Paris Perdikaris. Understanding and mitigating gradient pathologies in physics-informed neural networks. *SIAM Journal on Scientific Computing*, 43(5):A3055–A3081, 2021.

[24] Chengping Rao, Hao Sun, and Yang Liu. Physics-informed deep learning for incompressible laminar flows. *Theoretical and Applied Mechanics Letters*, 10(3):207–212, 2020.

[25] Shengze Cai, Zhiping Mao, Zhicheng Wang, Minglang Yin, and George Em Karniadakis. Physics-informed neural networks (pinns) for fluid mechanics: A review. *Acta Mechanica Sinica*, 37(12):1727–1738, 2021.

[26] Rafael Bischof and Michael Anton Kraus. Multi-objective loss balancing for physics-informed deep learning. *Computer Methods in Applied Mechanics and Engineering*, 439:117914, 2025.

[27] José X. Matos, C. Fernandes, and M. A. Alves. Simulation of steady viscoelastic creeping flows with physics-informed neural networks. *Journal of Non-Newtonian Fluid Mechanics*, 2026. in press.